

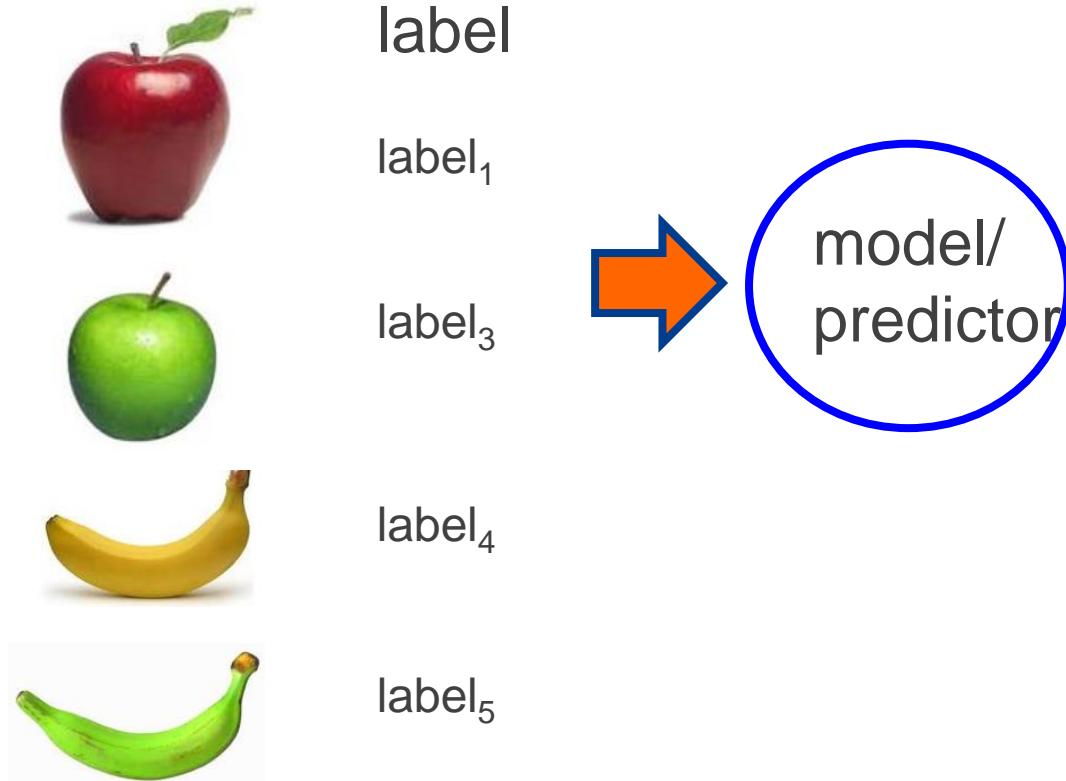
# COMP9321: Data services engineering

## Week 8: Clustering

Term3, 2019

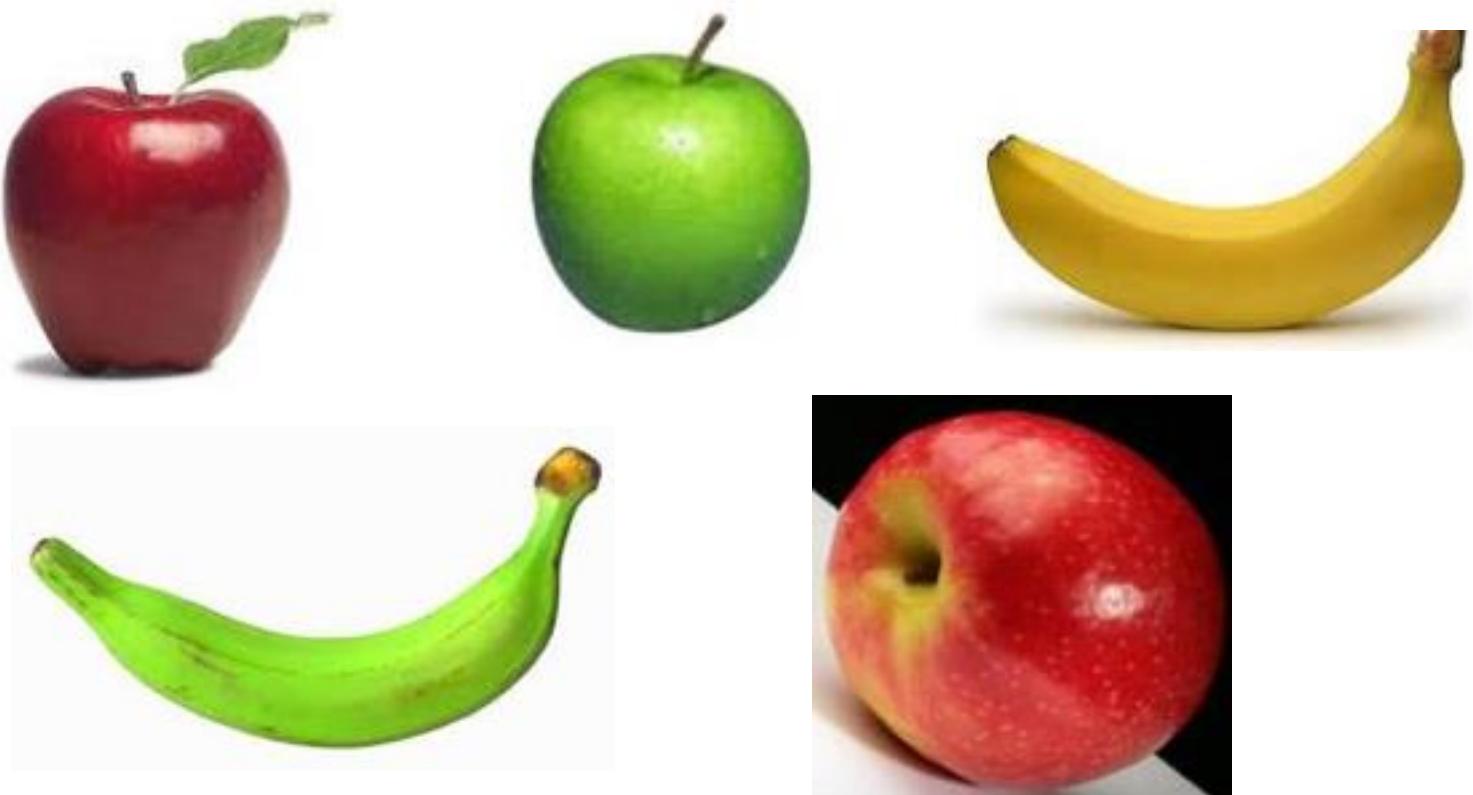
By Mortada Al-Banna, CSE UNSW

# Supervised learning



Supervised learning: given labeled examples

# Unsupervised learning



Unsupervised learning: given data, i.e. examples, but no labels

# Unsupervised Learning

Definition of Unsupervised Learning:

Learning useful structure *without* labeled classes, optimization criterion, feedback signal, or any other information beyond the raw data

# Unsupervised Learning

- Unsupervised learning involves operating on datasets without labelled responses or target values.
- The goal is to capture a structure of interest of useful information (e.g., relationships)
- Unsupervised learning good be used in:
  - Visualizing the structure of a complex dataset
  - Compressing and summarising the data (e.g, image compression)
  - Extracting features for supervised learning
  - Discover groups or outliers

# Clustering

## Unsupervised Learning

# Clustering

- Unsupervised learning
- Requires data, but no labels
- Detect patterns

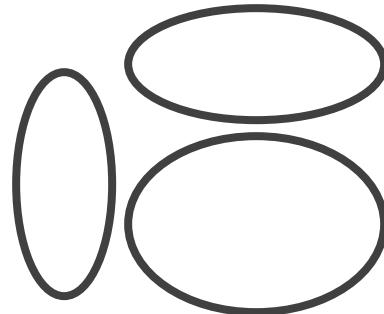
# Motivations of Clustering

- exploratory data analysis
  - understanding general characteristics of data
  - visualizing data
- generalization – infer something about an instance (e.g. a gene) based on how it relates to other instances

# Paradigms

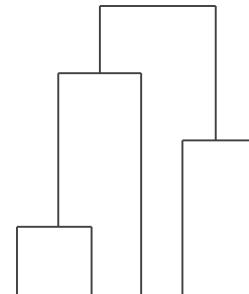
## Flat algorithms

- Usually start with a random (partial) partitioning
- Refine it iteratively
  - $K$  means clustering
  - Model based clustering
- Spectral clustering



## Hierarchical algorithms

- Bottom-up, agglomerative
- Top-down, divisive



# Paradigms

Hard clustering: Each example belongs to exactly one cluster

Soft clustering: An example can belong to more than one cluster (probabilistic)

- Makes more sense for applications like creating browsable hierarchies
- You may want to put a pair of sneakers in two clusters: (i) sports apparel and (ii) shoes

# Clustering: Image Segmentation

Break up the image into meaningful or perceptually similar regions



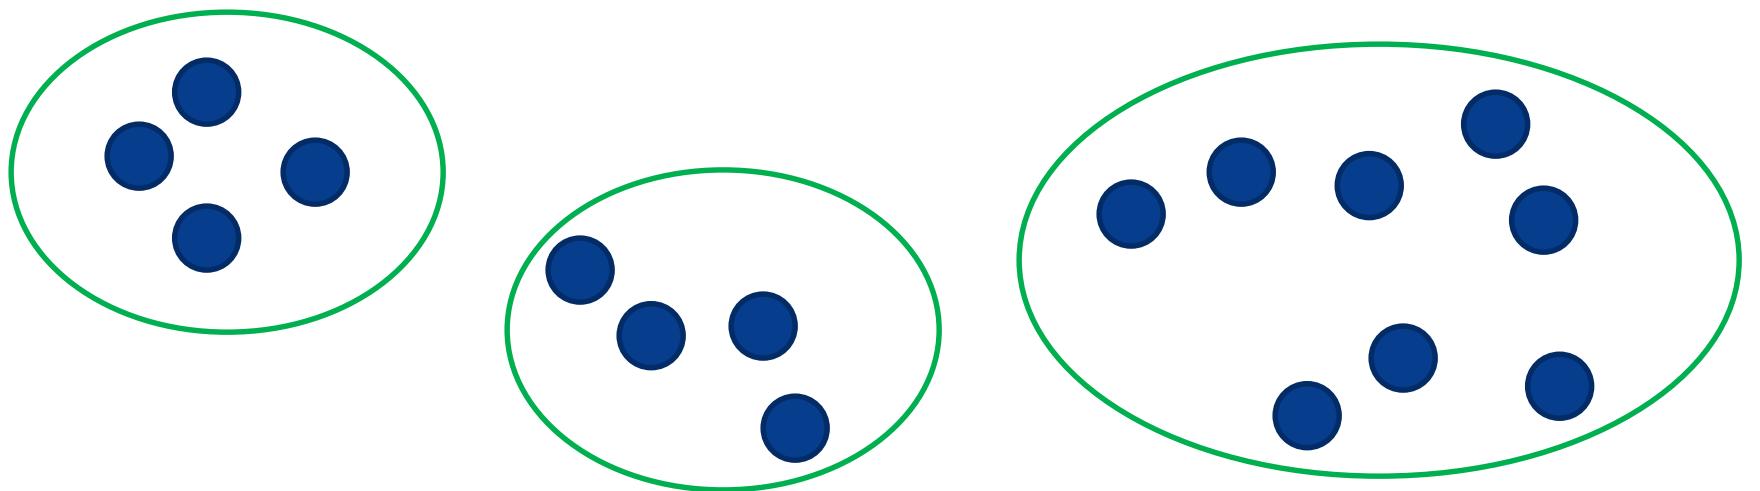
Figure from: James Hayes

# Clustering: Edge Detection

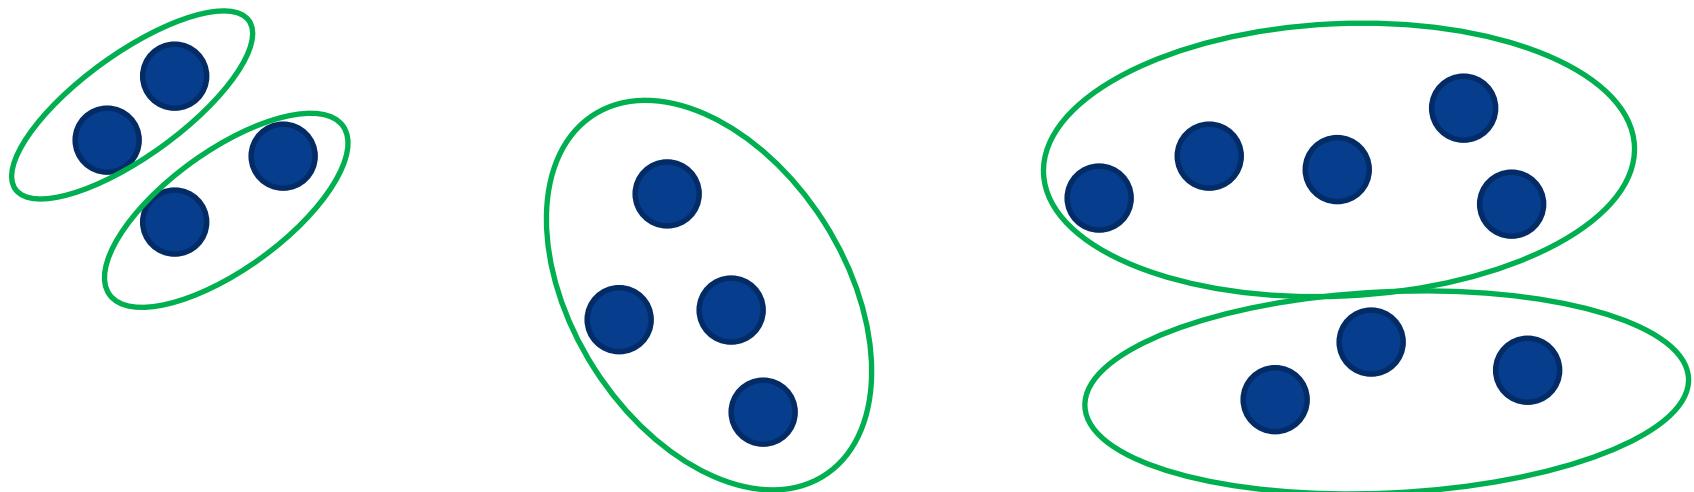


Figure from: James Hayes

# Basic Idea of Clustering



# Basic Idea of Clustering



# Basic Idea of Clustering

Group together similar data points (instances)

- How to measure the similarity?
- ✓ What could similar mean?
- How many clusters do we need?

# K-means

Most well-known and popular clustering algorithm:

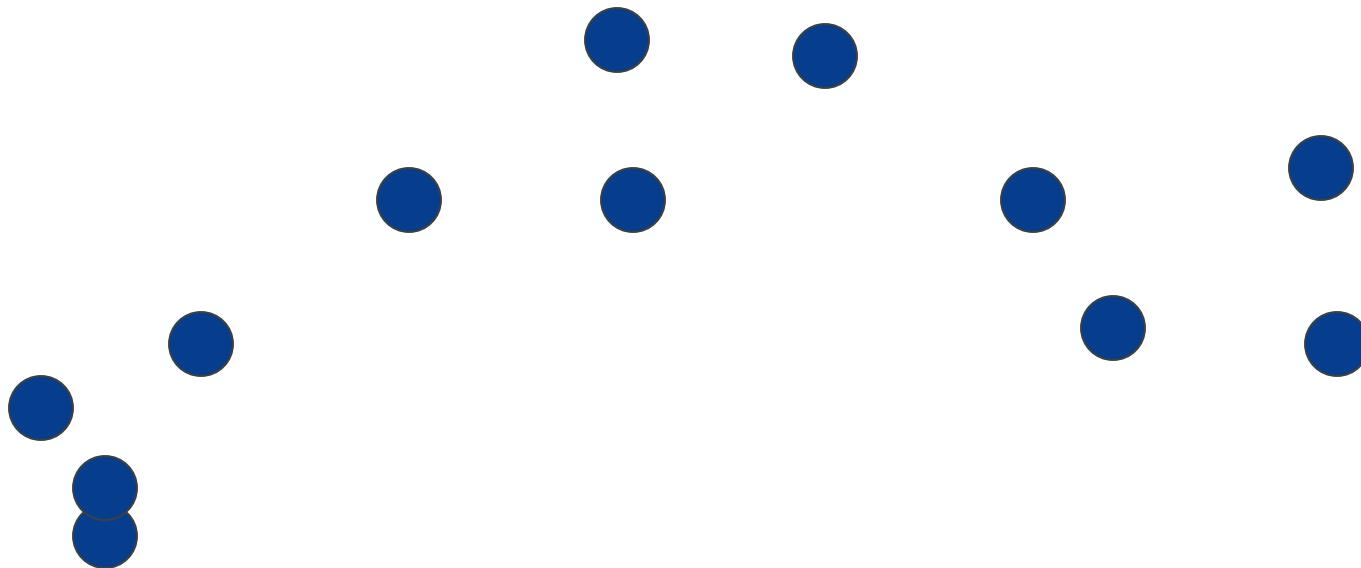
Step 1. Start with some initial cluster centers ( $k$  random points)

Step 2. Iterate:

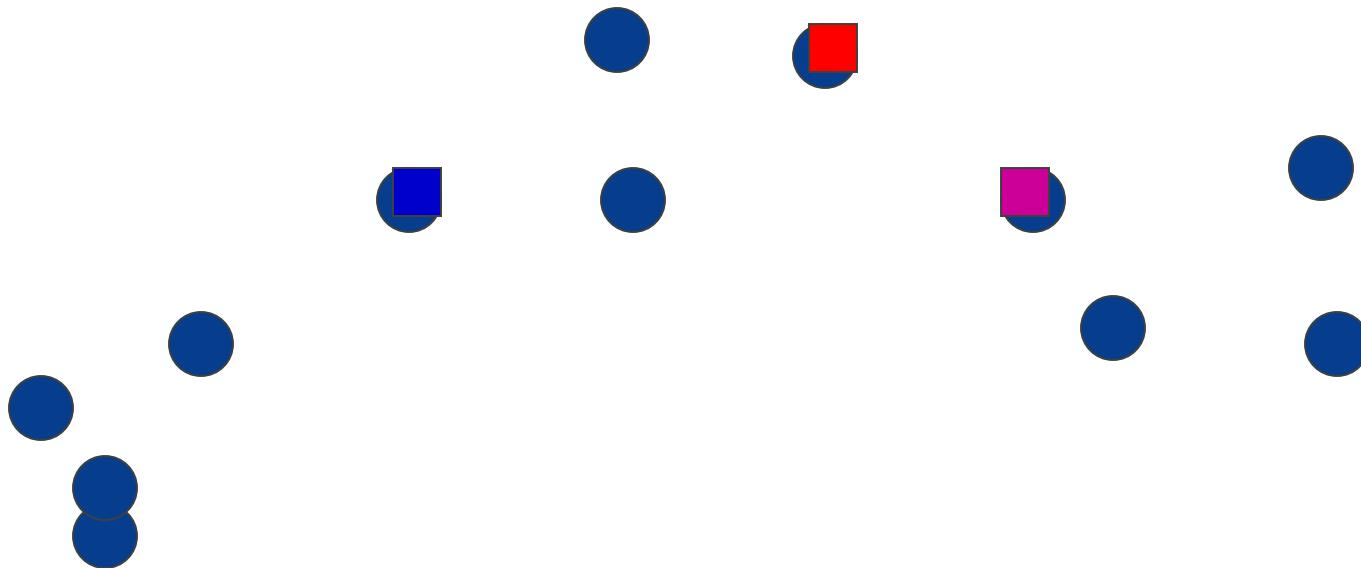
- Assign/cluster each example to closest center
- Recalculate and change centers as the mean of the points in the cluster.

Step 3. Stop when no points' assignments change

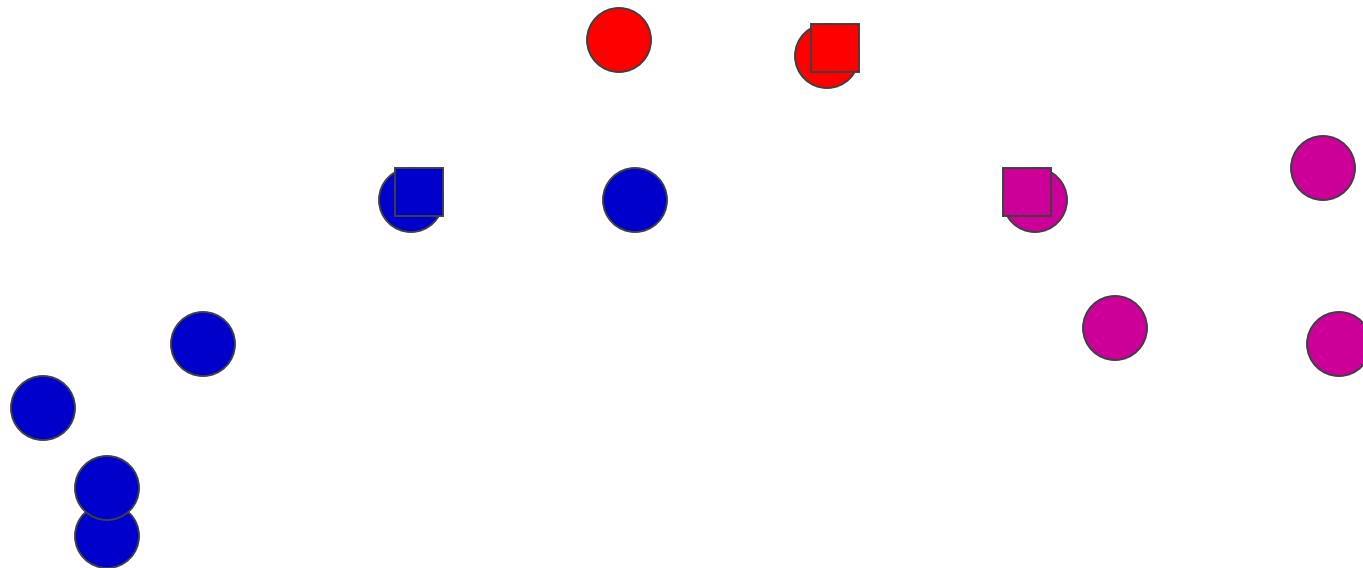
# K-means: an example



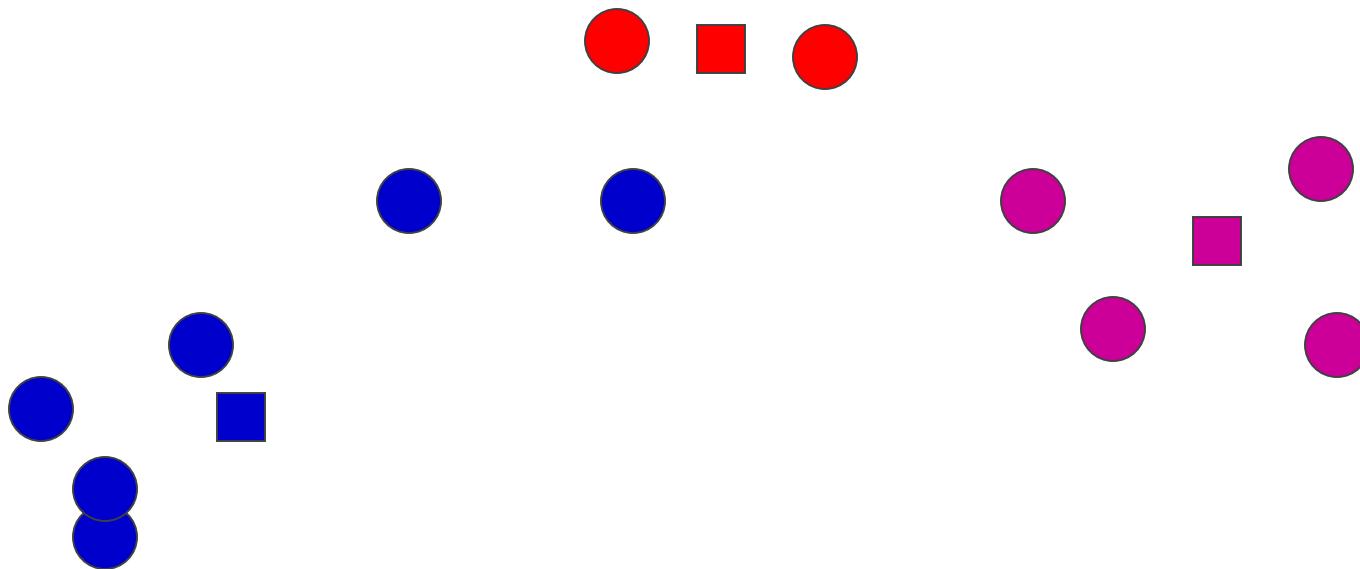
# K-means: Initialize centers randomly



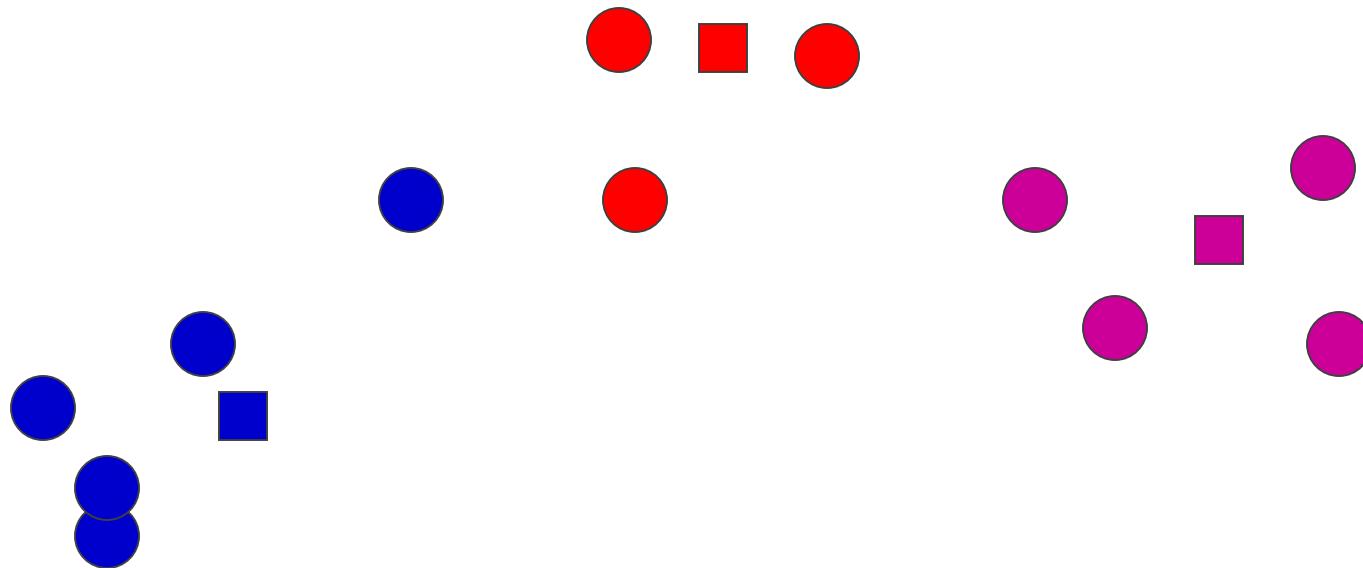
# K-means: assign points to nearest center



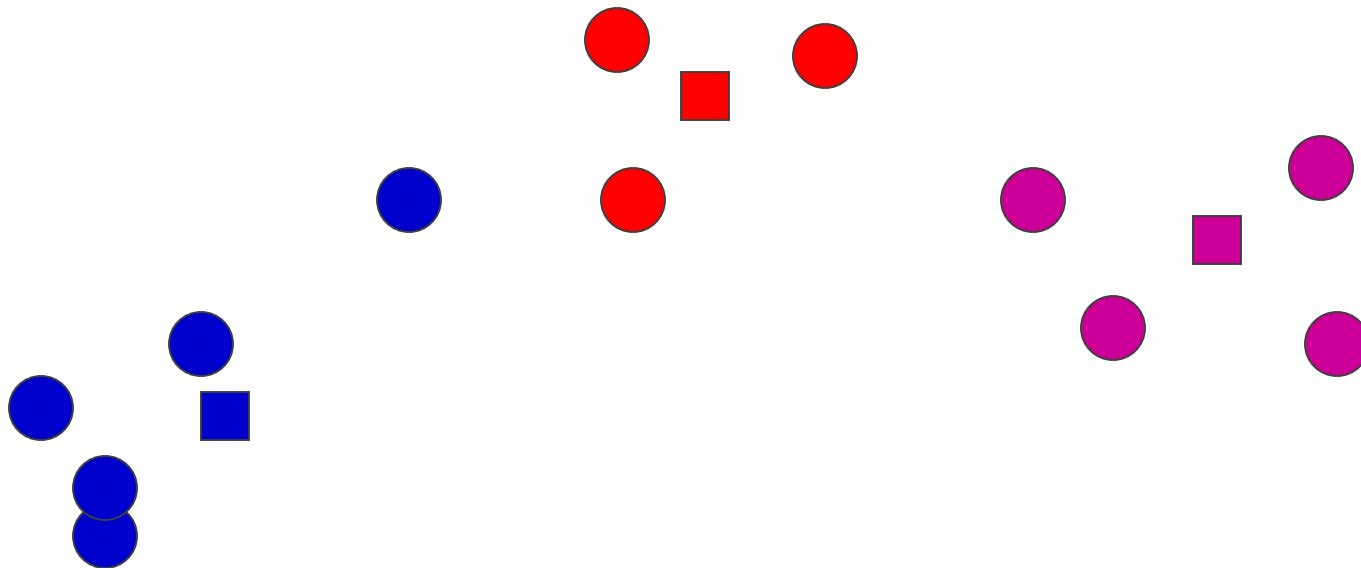
# K-means: readjust centers



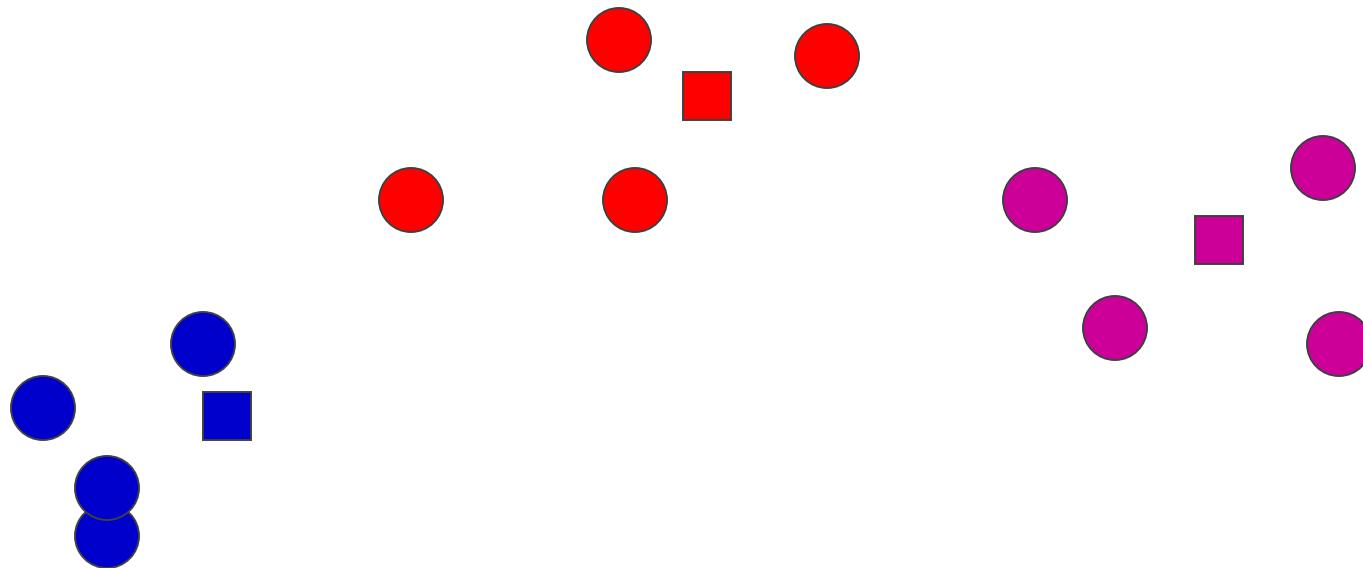
# K-means: assign points to nearest center



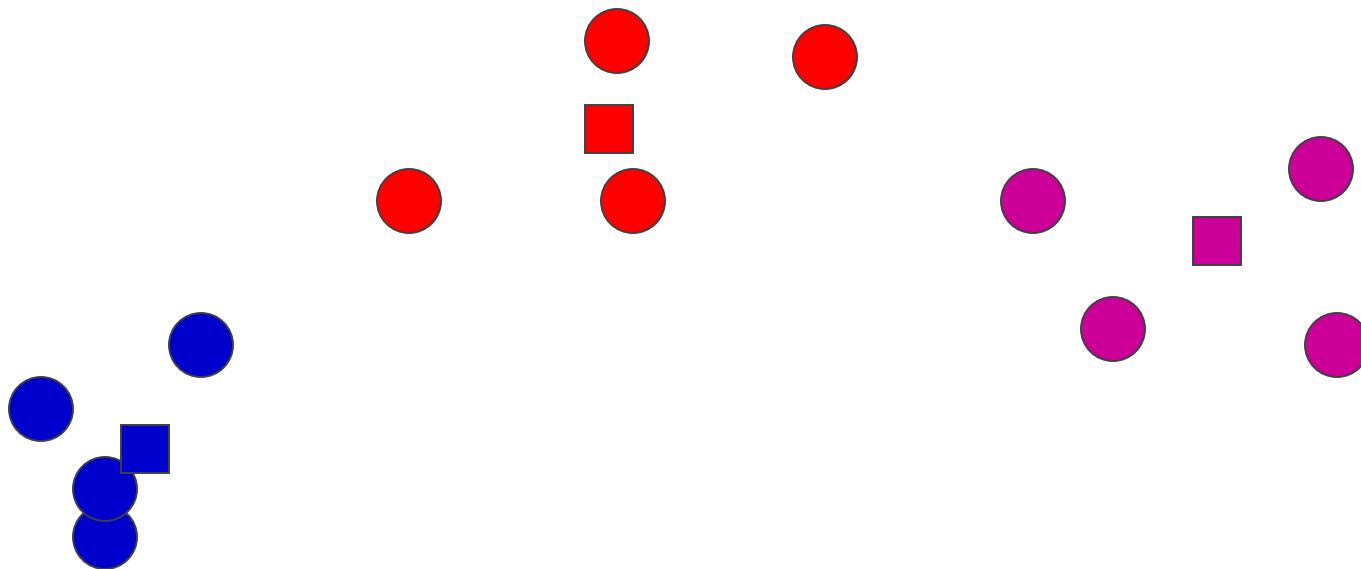
# K-means: readjust centers



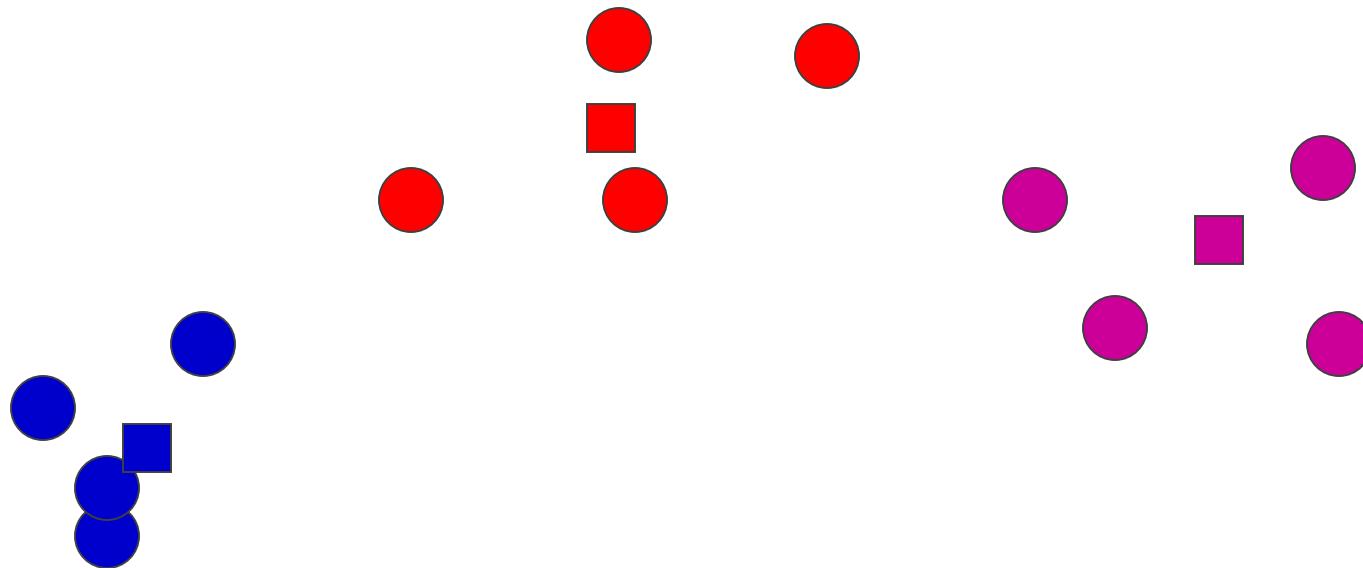
# K-means: assign points to nearest center



# K-means: readjust centers



# K-means: assign points to nearest center

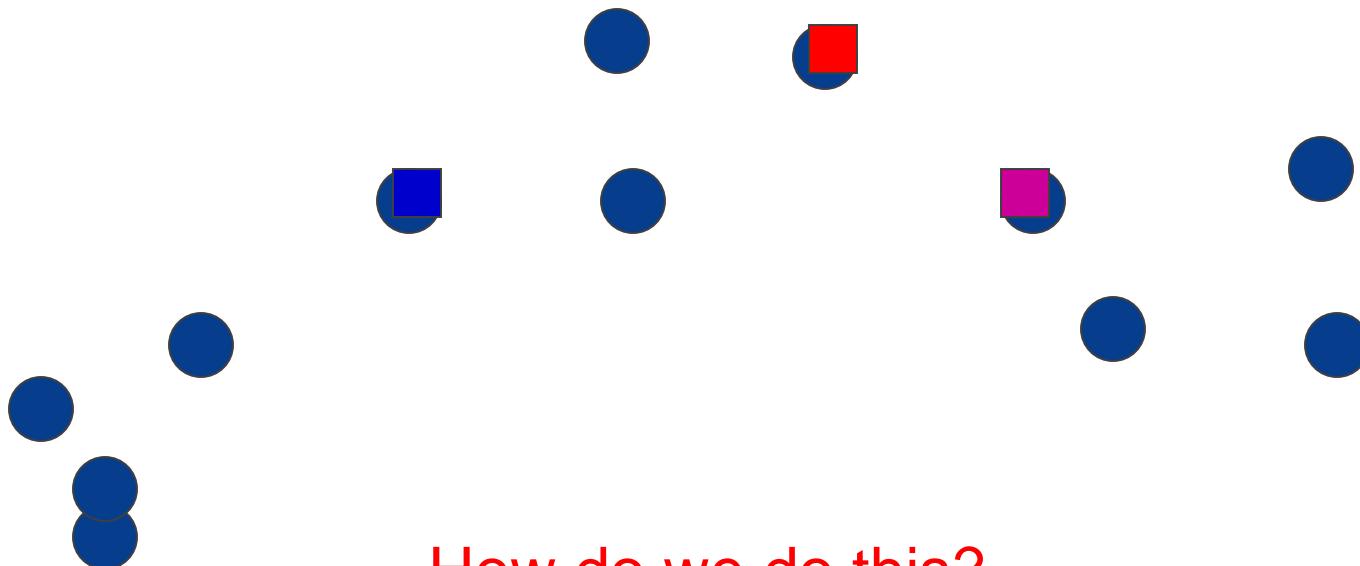


No changes: Done

# K-means

Iterate:

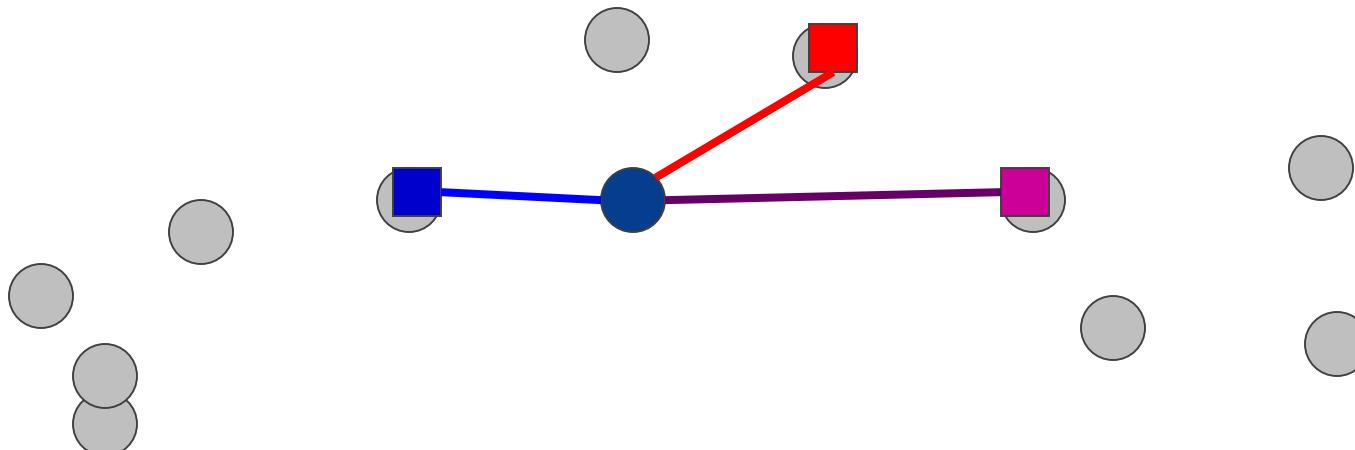
- **Assign/cluster each example to closest center**
- Recalculate centers as the mean of the points in a cluster



# K-means

Iterate:

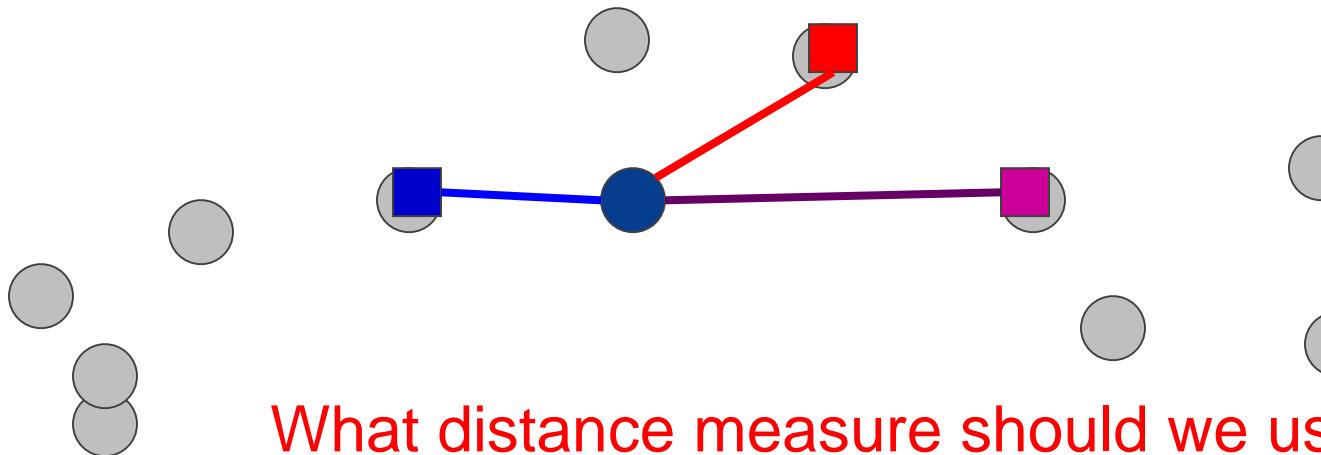
- **Assign/cluster each example to closest center**
  - iterate over each point:
    - get distance to each cluster center
    - assign to closest center (hard cluster)
- Recalculate centers as the mean of the points in a cluster



# K-means

Iterate:

- **Assign/cluster each example to closest center**
  - iterate over each point:
    - get **distance** to each cluster center
    - assign to closest center (hard cluster)
- Recalculate centers as the mean of the points in a cluster



What distance measure should we use?

# Distance measures

Euclidean:

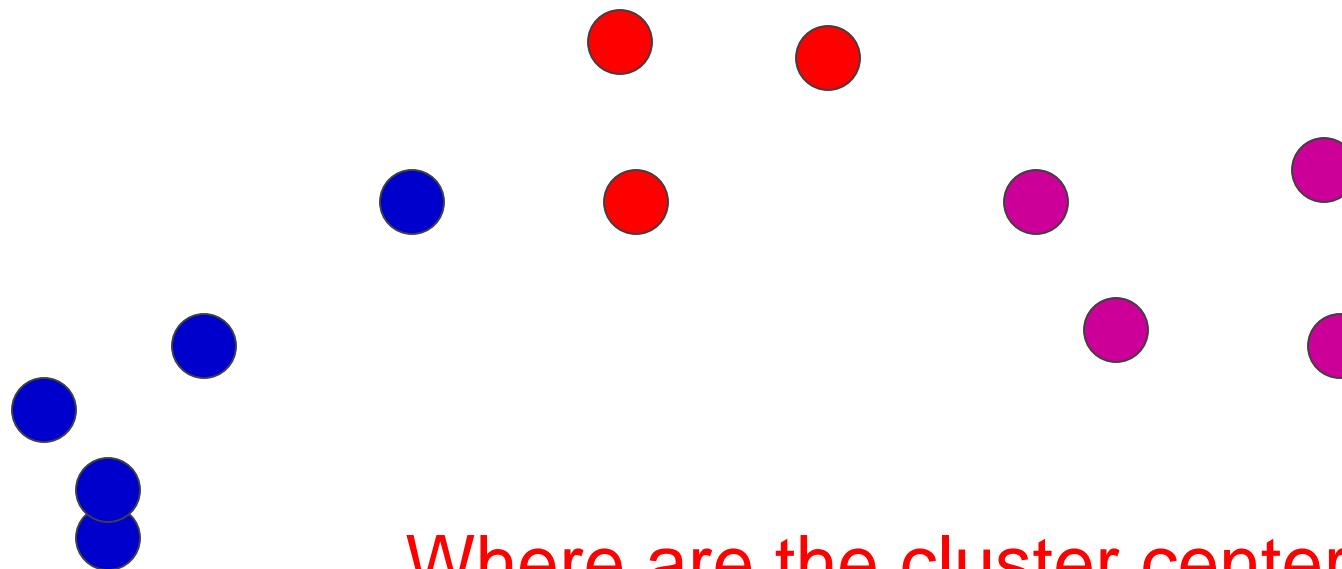
$$d(x, y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

good for spatial data

# K-means

Iterate:

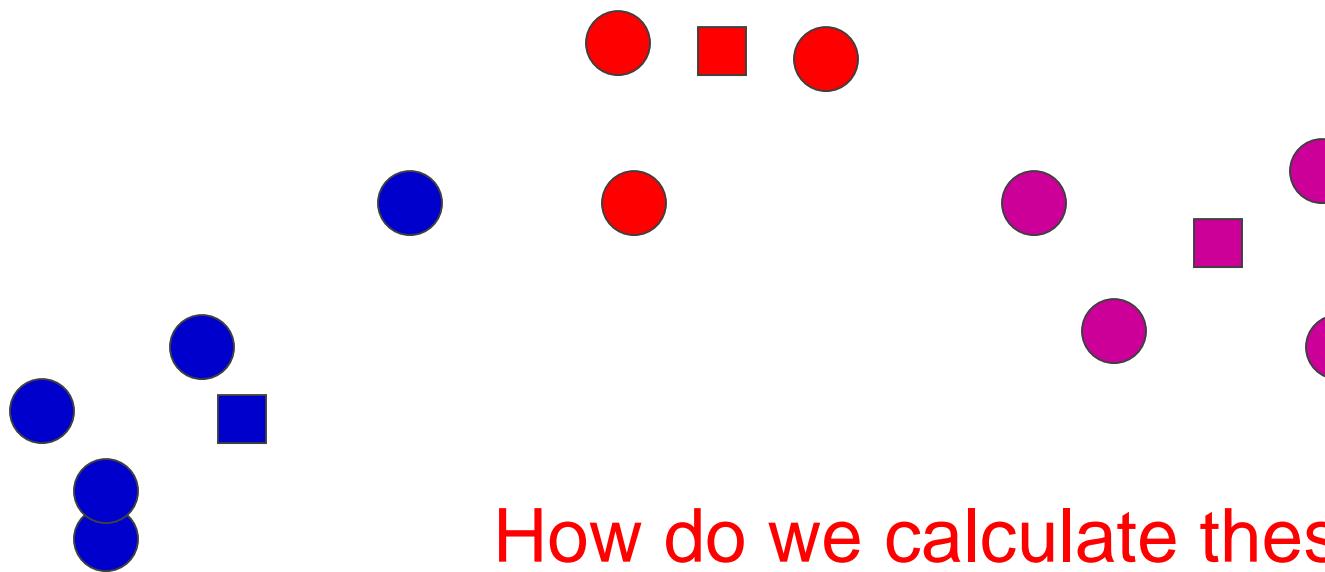
- Assign/cluster each example to closest center
- Recalculate centers as the mean of the points in a cluster



# K-means

Iterate:

- Assign/cluster each example to closest center
- Recalculate centers as the mean of the points in a cluster

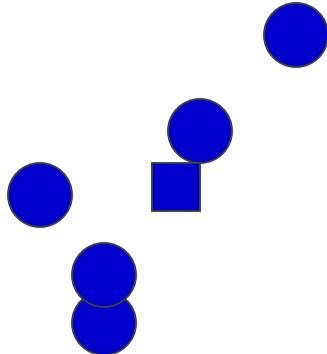


# K-means

Iterate:

- Assign/cluster each example to closest center
- Recalculate centers as the mean of the points in a cluster

e.g., for a set of instances that have been assigned to a cluster  $\mathcal{C}_j$ , we compute the mean of the cluster as follow



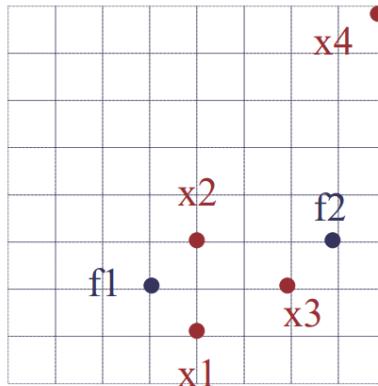
$$\mu(c_j) = \frac{\sum_{\vec{x}_i \in c_j} \vec{x}_i}{|c_j|}$$

# K-means

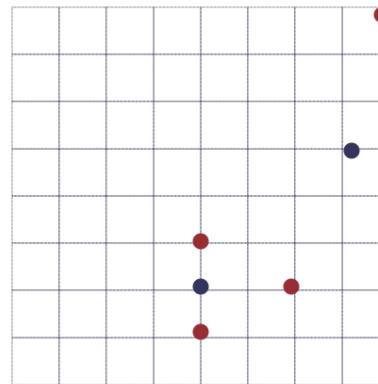
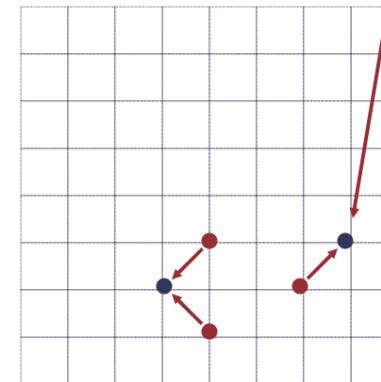
```
given : a set  $X = \{\vec{x}_1 \dots \vec{x}_n\}$  of instances  
select  $k$  initial cluster centers  $\vec{f}_1 \dots \vec{f}_k$   
while stopping criterion not true do  
    for all clusters  $c_j$  do  
        // determine which instances are assigned to this cluster  
         $c_j = \left\{ \vec{x}_i \mid \forall f_l \text{ dist}(\vec{x}_i, \vec{f}_j) < \text{dist}(\vec{x}_i, \vec{f}_l) \right\}$   
    for all means  $\vec{f}_j$  do  
        // update the cluster center  
         $\vec{f}_j = \mu(c_j)$ 
```

# Run an example together ~~

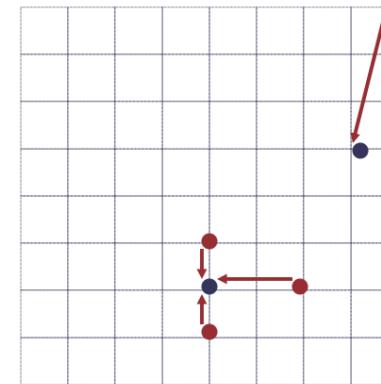
Initialization: 4 points, 2 clusters and distance function



$$dist(x_i, x_j) = \sum_e |x_{i,e} - x_{j,e}|$$
$$f_1 = \left\langle \frac{4+4}{2}, \frac{1+3}{2} \right\rangle = \langle 4, 2 \rangle$$
$$f_2 = \left\langle \frac{6+8}{2}, \frac{2+8}{2} \right\rangle = \langle 7, 5 \rangle$$



$$f_1 = \left\langle \frac{4+4+6}{3}, \frac{1+3+2}{3} \right\rangle = \langle 4.67, 2 \rangle$$
$$f_2 = \left\langle \frac{8}{1}, \frac{8}{1} \right\rangle = \langle 8, 8 \rangle$$



# Properties of K-means

Guaranteed to converge in a finite number of iterations

Running time per iteration

1. Assign data points to closest cluster center  
 $O(KN)$  time
2. Change the cluster center to the average of its assigned points  $O(N)$

# K-means variations/parameters

Start with some initial cluster centers

Iterate:

- Assign/cluster each example to closest center
- Recalculate centers as the mean of the points in a cluster

What are some other  
variations/parameters we haven't  
specified?

# K-means variations/parameters

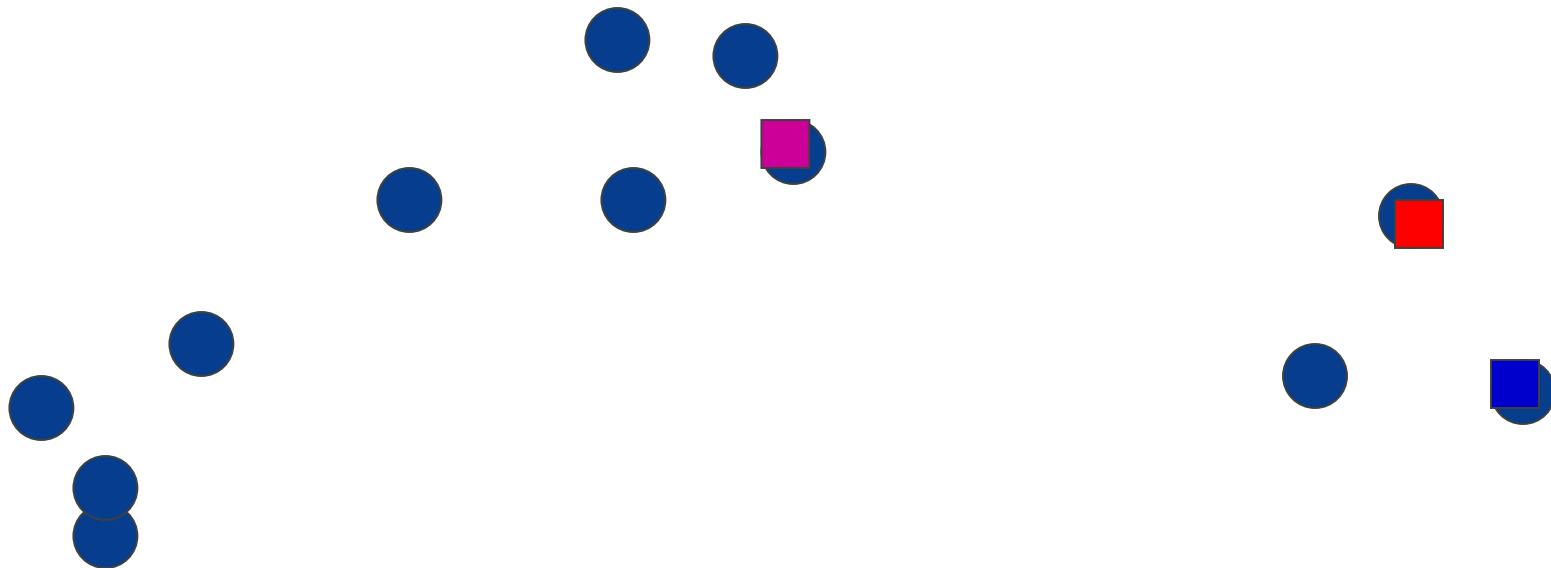
Initial (seed) cluster centers

Convergence

- A fixed number of iterations
- partitions unchanged
- Cluster centers don't change

K!

# K-means: Initialize centers randomly



What would happen here?

Seed selection ideas?

# Seed choice

Results can vary drastically based on random seed selection

Some seeds can result in poor convergence rate, or convergence to sub-optimal clustering

## Common heuristics

- Random centers in the space
- Randomly pick examples
- Points least similar to any existing center (furthest centers heuristic)
- **Try out multiple starting points**
- Initialize with the results of another clustering method

# Furthest centers heuristic

$\mu_1$  = pick random point

for  $i = 2$  to  $K$ :

$\mu_i$  = point that is furthest from **any** previous centers

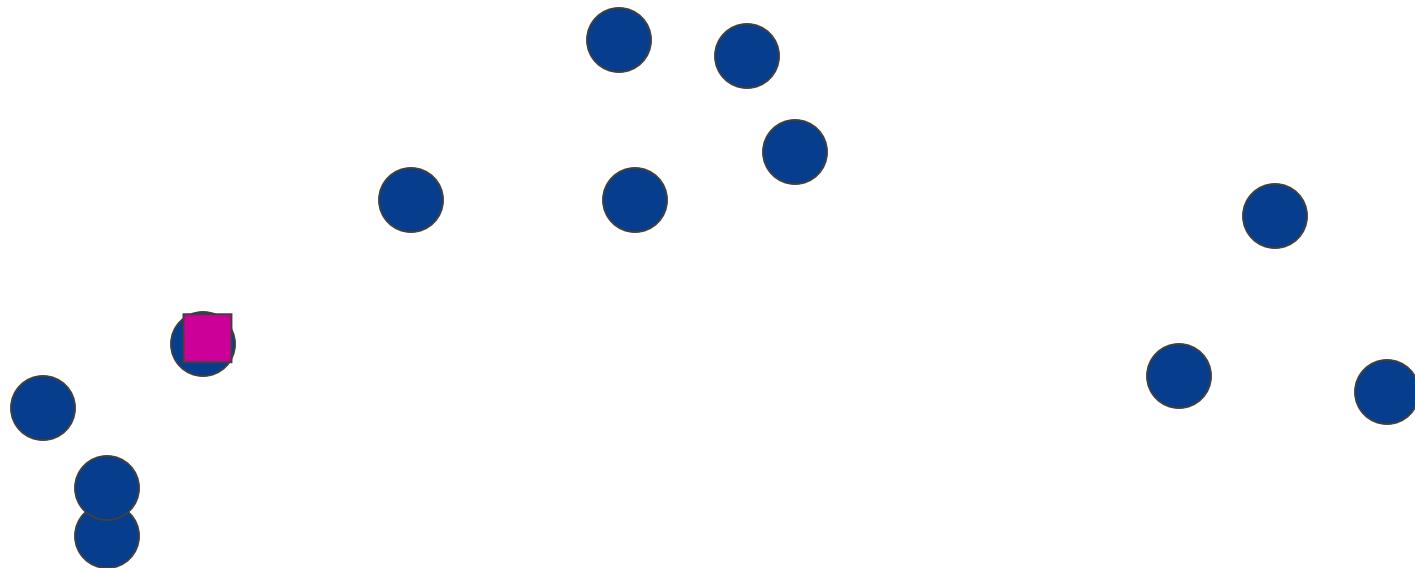
---

$$m_i = \arg \max_x \min_{m_j : 1 < j < i} d(x, m_j)$$

point with the largest  
distance to any previous  
center

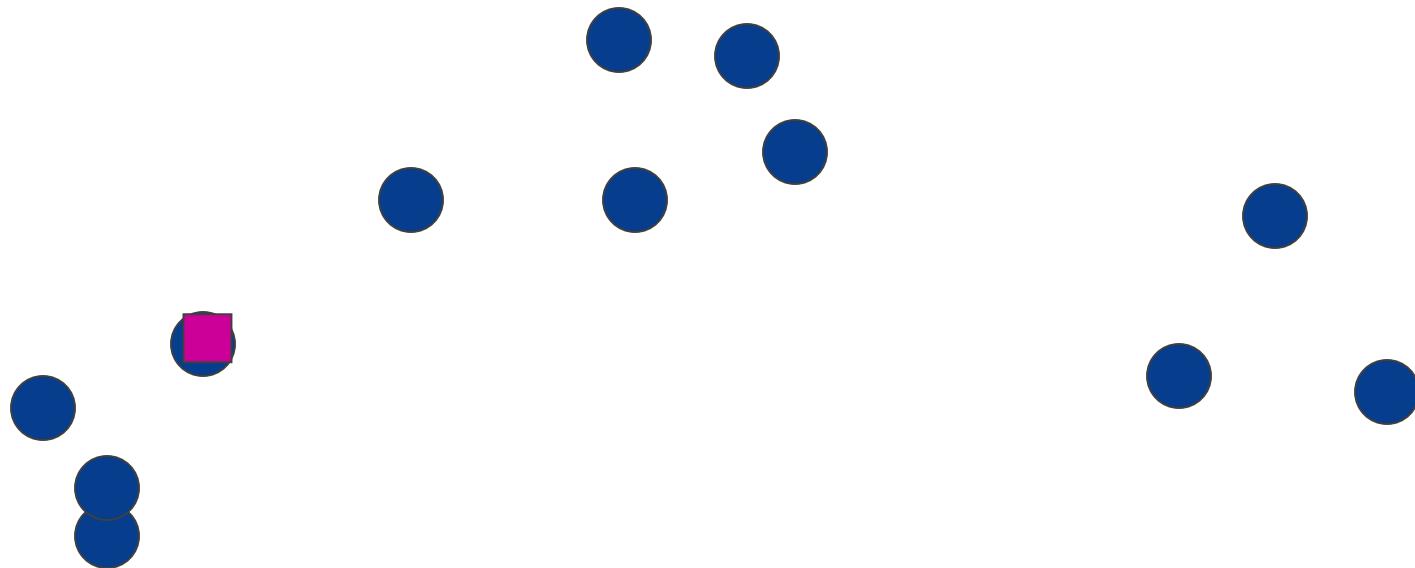
smallest distance from  $x$  to  
any previous center

# K-means: Initialize furthest from centers



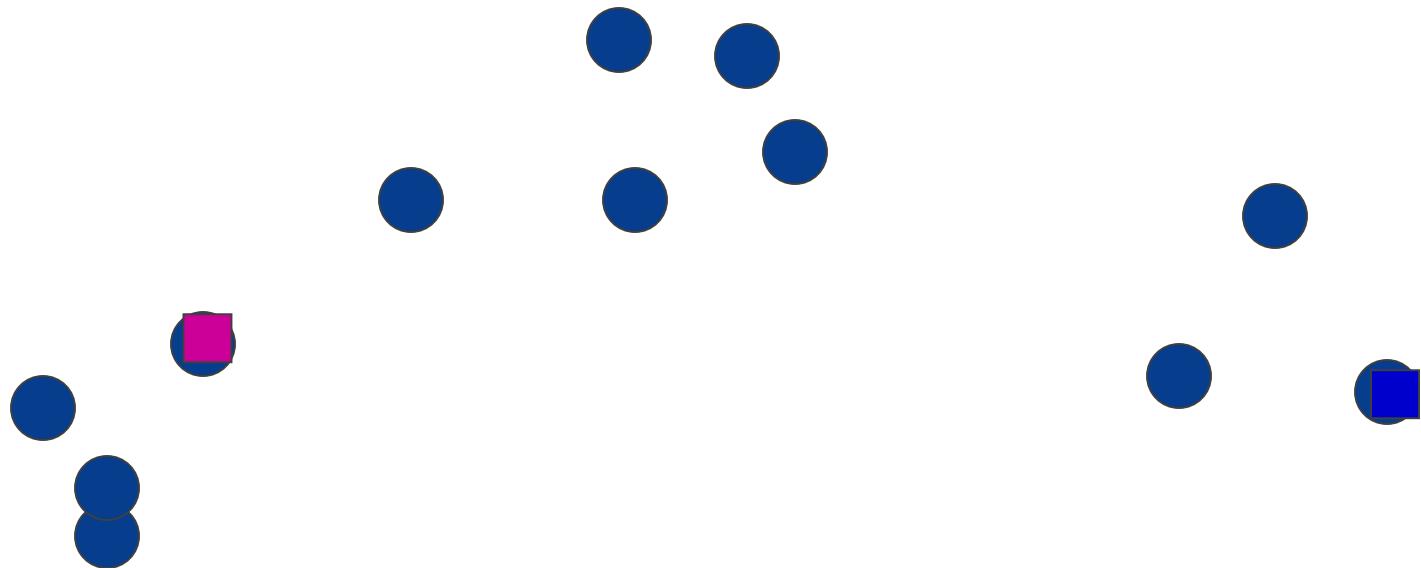
Pick a random point for the first center

# K-means: Initialize furthest from centers



What point will be chosen next?

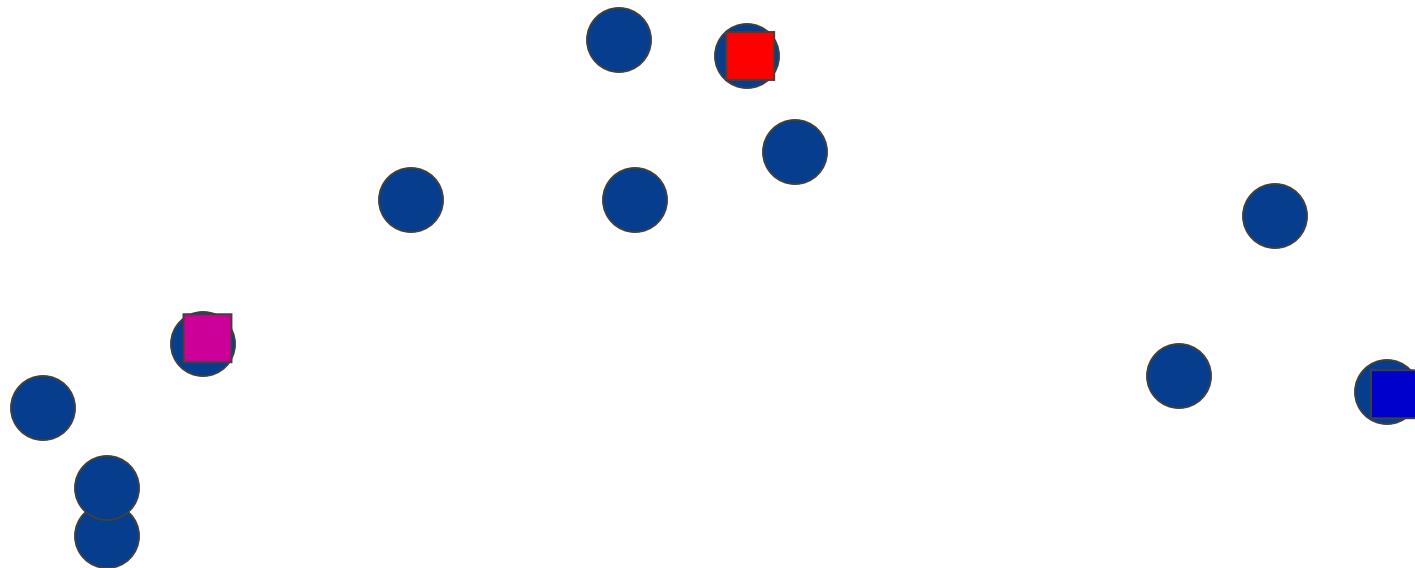
# K-means: Initialize furthest from centers



Furthest point from center

What point will be chosen next?

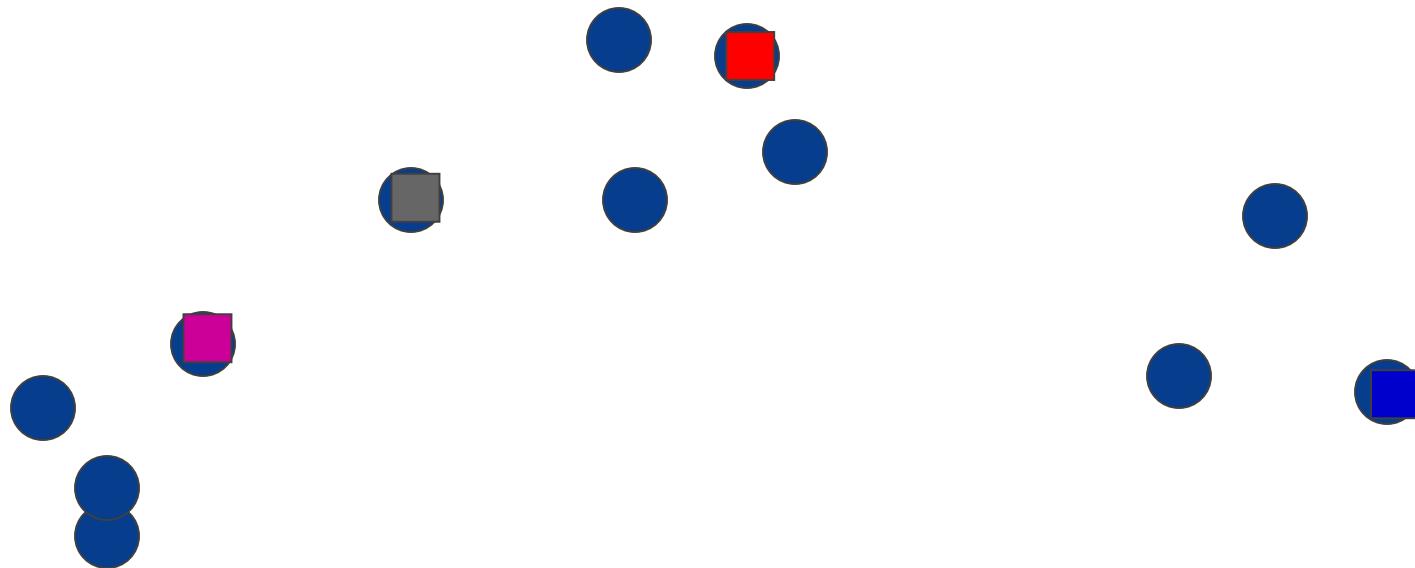
# K-means: Initialize furthest from centers



Furthest point from center

What point will be chosen next?

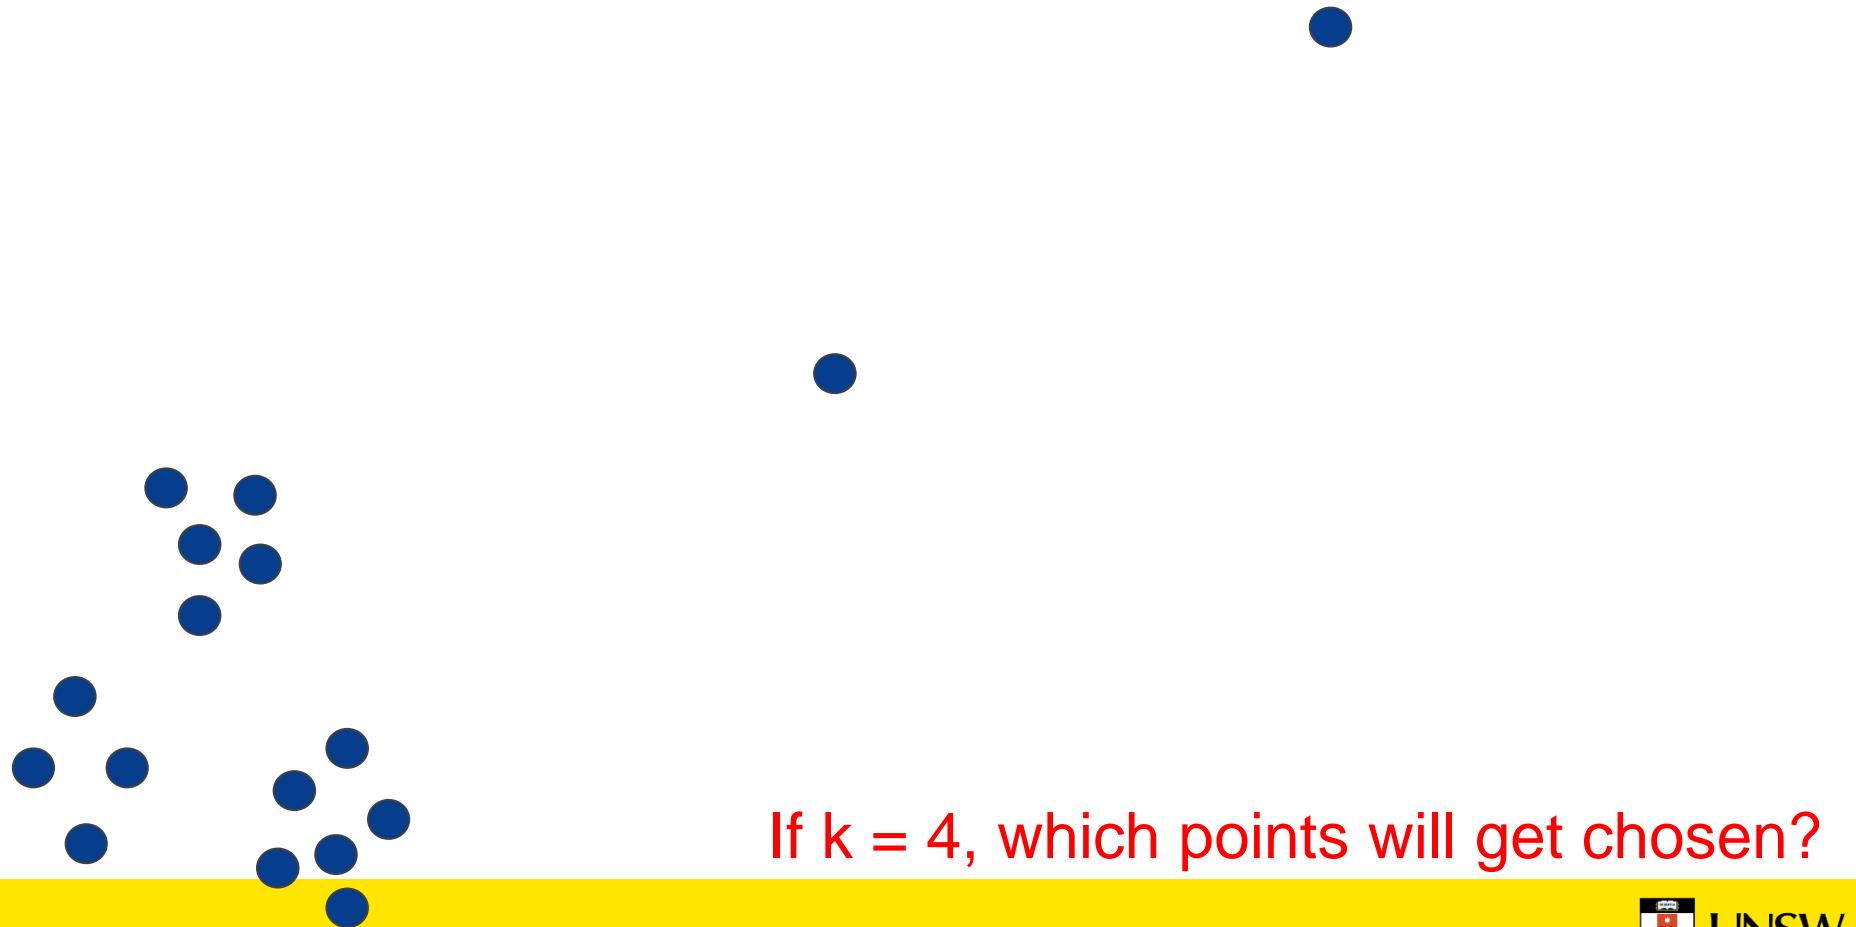
# K-means: Initialize furthest from centers



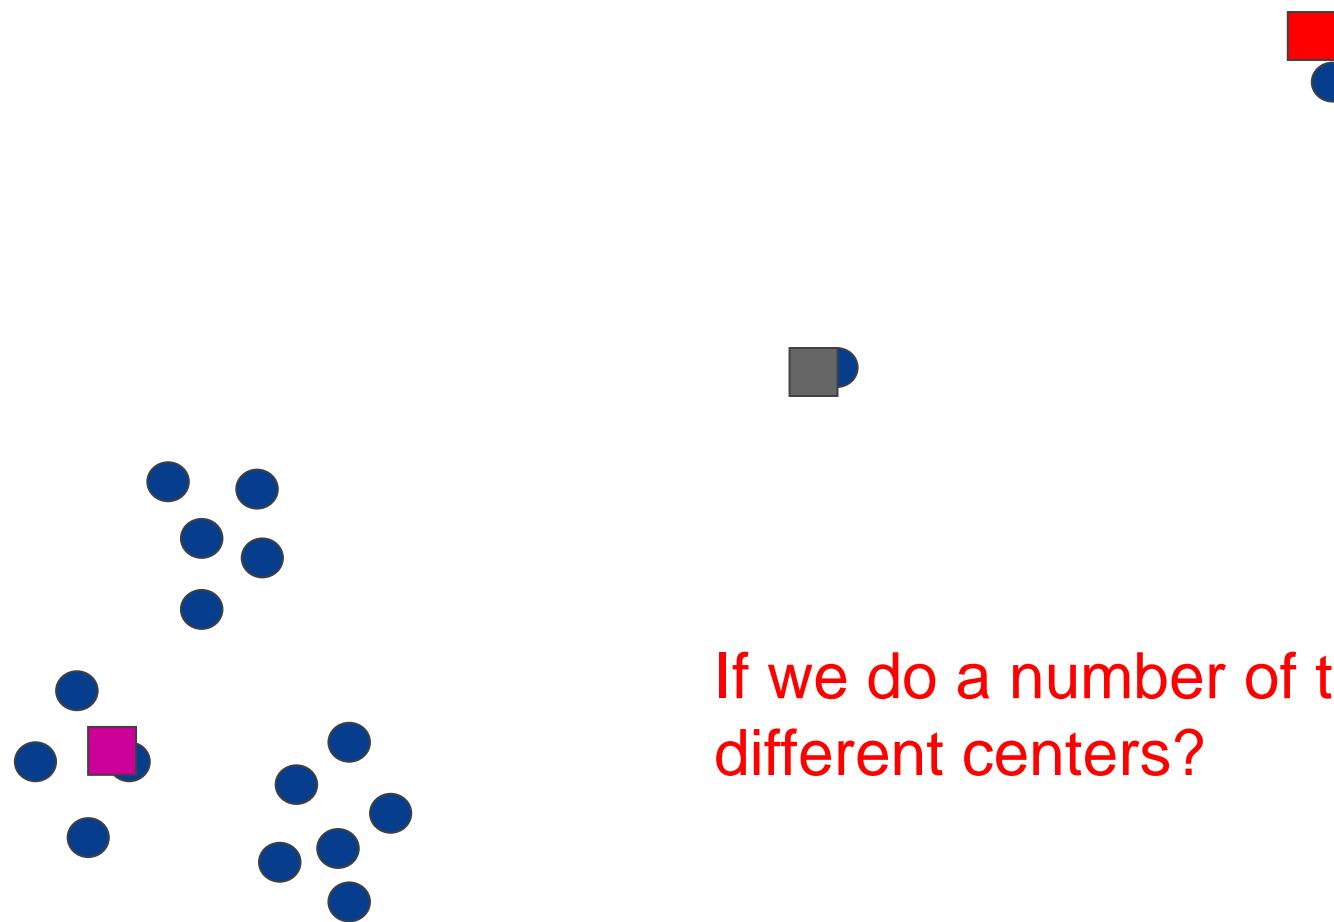
Furthest point from center

Any issues/concerns with this approach?

# Furthest points concerns



# Furthest points concerns



If we do a number of trials, will we get different centers?

# K-means++

$\mu_1$  = pick random point

for  $k = 2$  to  $K$ :

    for  $i = 1$  to  $N$ :

$s_i = \min d(x_i, \mu_1 \dots \mu_{k-1})$  // smallest distance to any center

$\mu_k$  = randomly pick point ***proportionate*** to ***s***

How does this help?

# K-means++

$\mu_1$  = pick random point

for  $k = 2$  to  $K$ :

    for  $i = 1$  to  $N$ :

$s_i = \min d(x_i, \mu_1 \dots \mu_{k-1})$  // smallest distance to any center

$\mu_k$  = randomly pick point ***proportionate*** to ***s***

- Makes it possible to select other points
  - if #points >> #outliers, we will pick good points
- Makes it non-deterministic, which will help with random runs
- Nice theoretical guarantees!

# What Is A Good Clustering?

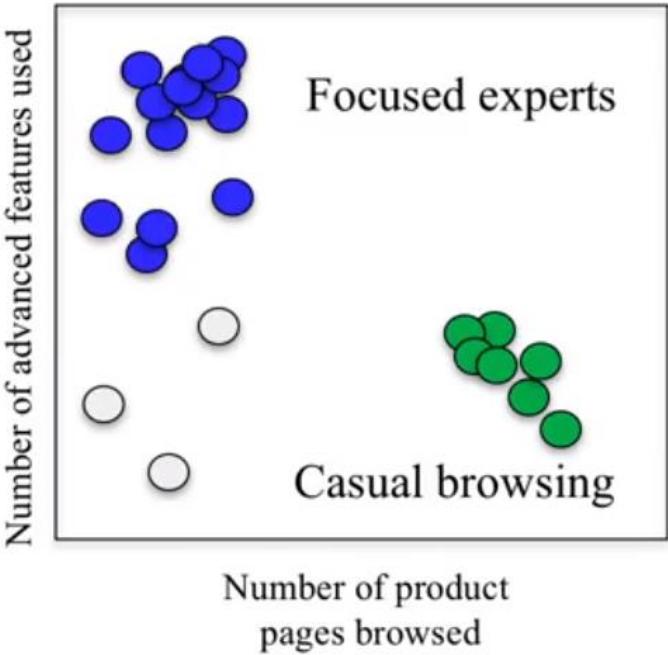
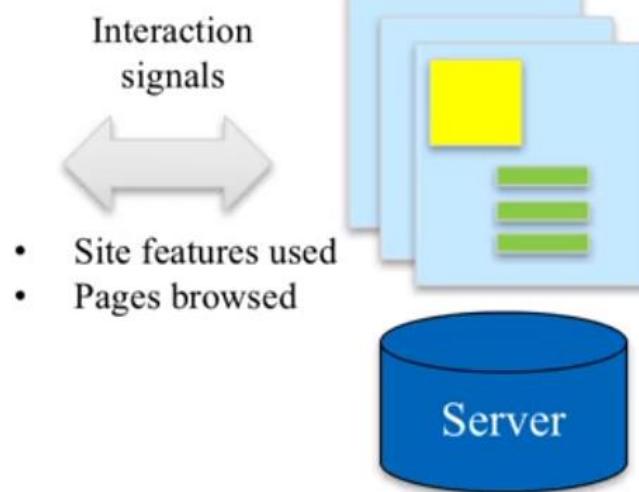
Internal criterion: A good clustering will produce high quality clusters in which:

- the intra-class (that is, intra-cluster) similarity is high
- the inter-class similarity is low
- The measured quality of a clustering depends on both the document representation and the similarity measure used

# Clustering Evaluation

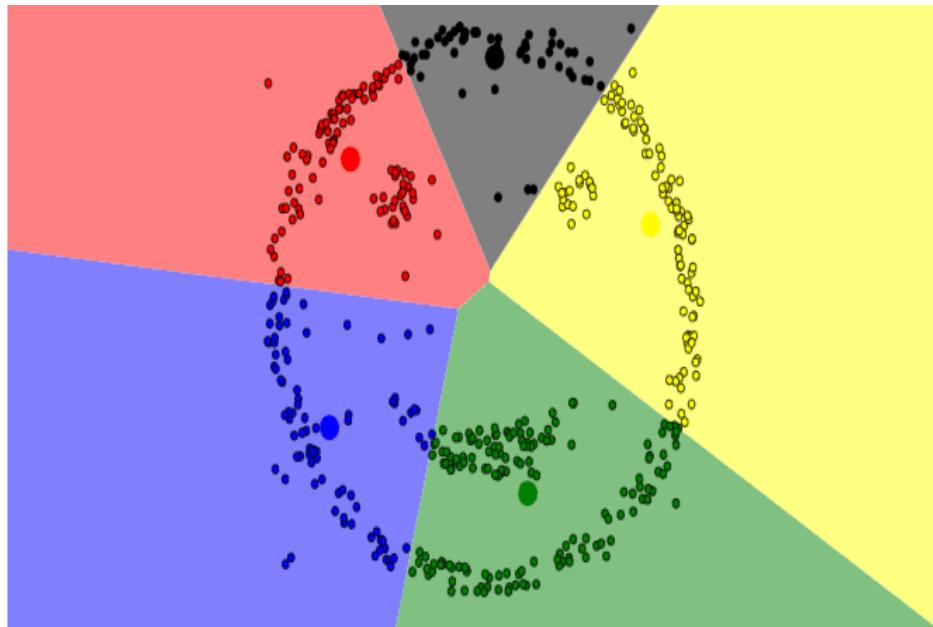
- Intra-cluster cohesion (compactness):
  - Cohesion measures how near the data points in a cluster are to the cluster centroid.
  - Sum of squared error (SSE) is a commonly used measure.
- Inter-cluster separation (isolation):
  - Separation means that different cluster centroids should be far away from one another.
- In most applications, expert judgments are still the key

# Web Clustering Examples



# Limitations of k-means

- Sometime the number of clusters is difficult to determine
- Does not do well with irregular or complex clusters.
- Has a problem with data containing outliers



# Q & A