

# COMP9321: Data services engineering

## Week 10: Introduction to Deep Learning

Term 3, 2019

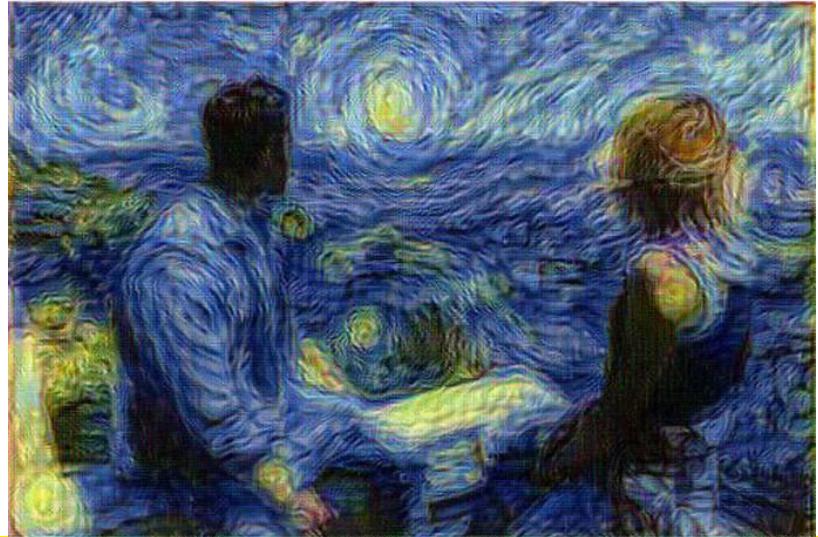
By Mortada Al-Banna, CSE UNSW





Cat

Dog



UNSW  
SYDNEY

# Why Deep Learning

- Deal very well with unstructured Data
- More Data = Better Learning
- More Accurate in some cases
- Extensive feature engineering is not required

# When to use Deep Learning

- Deep Learning outperform other techniques if the *data size is large*. But with small data size, traditional Machine Learning algorithms are preferable.
- Deep Learning techniques need to have *high end infrastructure* to train in reasonable time.
- When there is *lack of domain understanding for feature introspection*, Deep Learning techniques outshines others as you have to worry less about feature engineering.
- Deep Learning really shines when it comes to *complex problems such as image classification, natural language processing, and speech recognition*.

Many machine learning methods inspired by biology, e.g., the (human) brain  
Our brain has  $\sim 10^{11}$  neurons, each of which communicates (is connected)  
to  $\sim 10^4$  other neurons

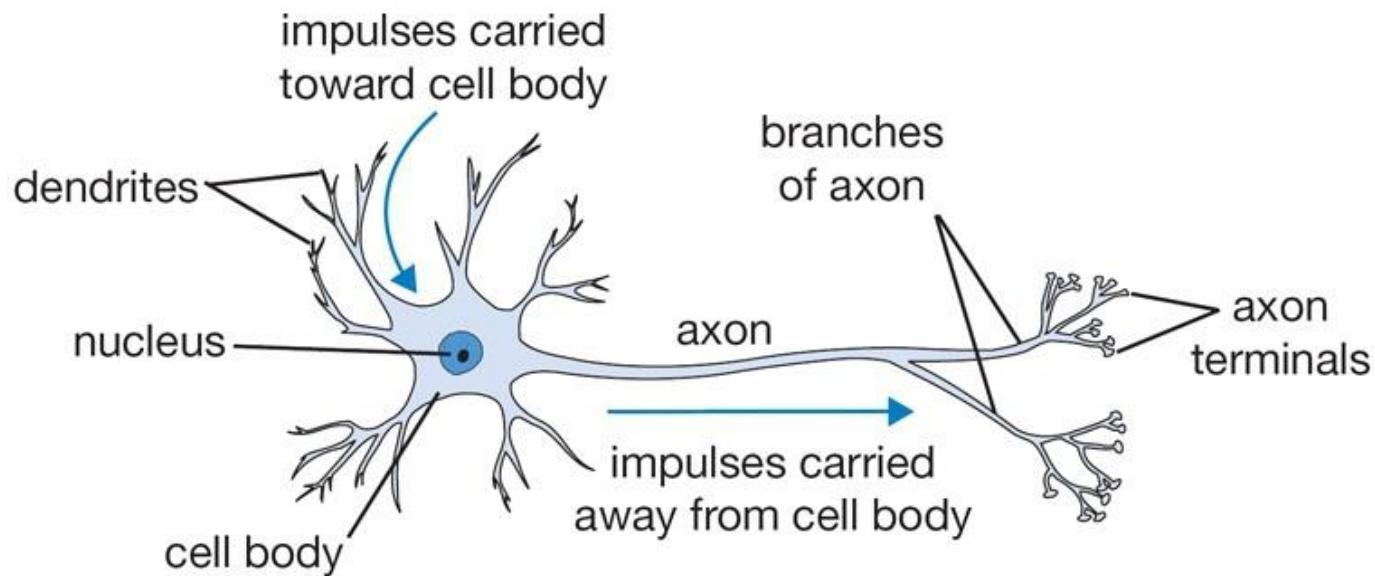


Figure : The basic computational unit of the brain: Neuron

[Pic credit: <http://cs231n.github.io/neural-networks-1/>]

Neural networks define functions of the inputs (**hidden features**), computed by neurons

Artificial neurons are called **units**

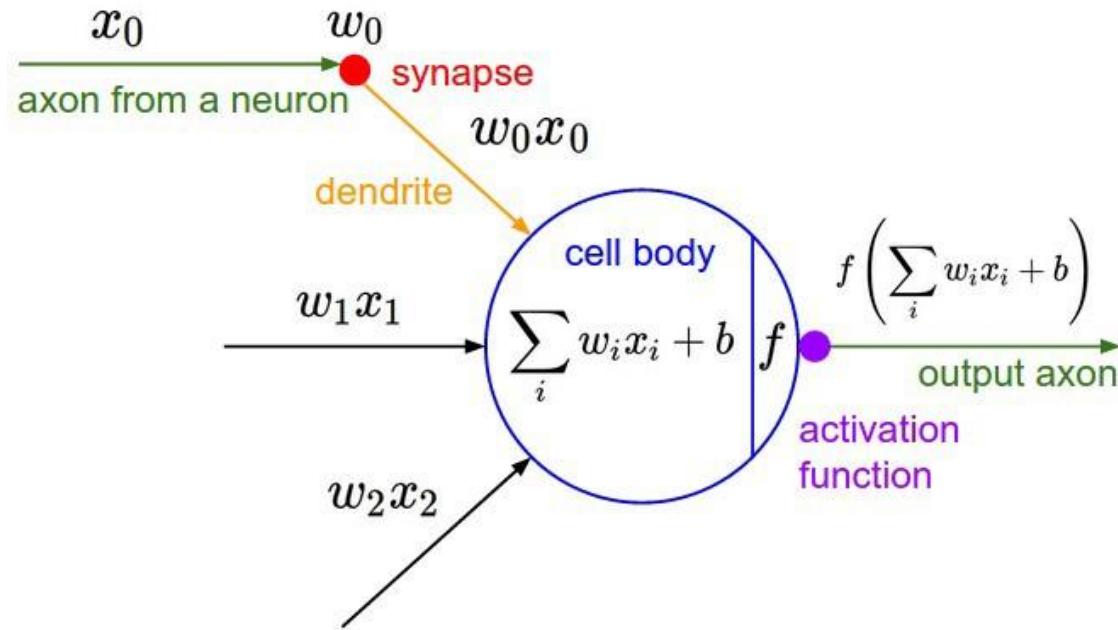


Figure : A mathematical model of the neuron in a neural network

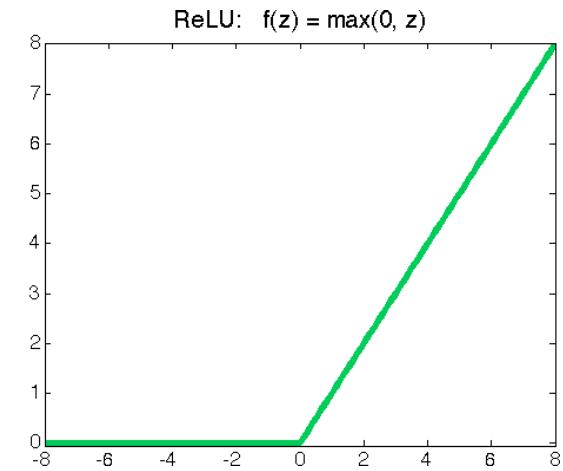
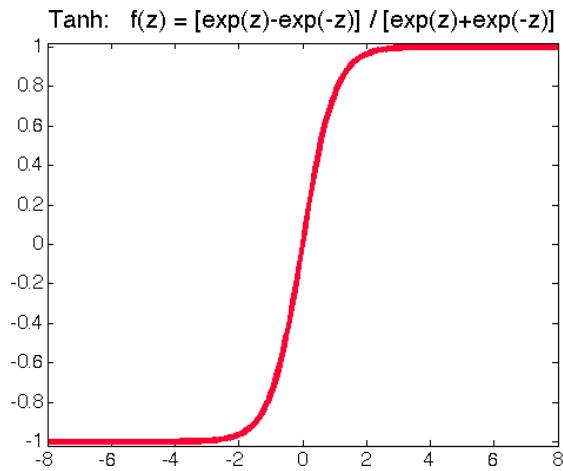
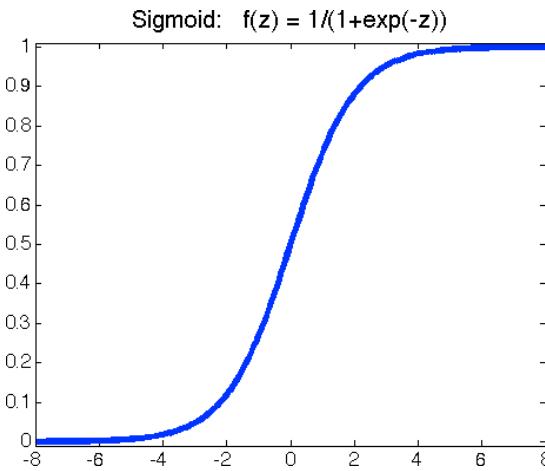
[Pic credit: <http://cs231n.github.io/neural-networks-1/>]

## Most commonly used activation functions:

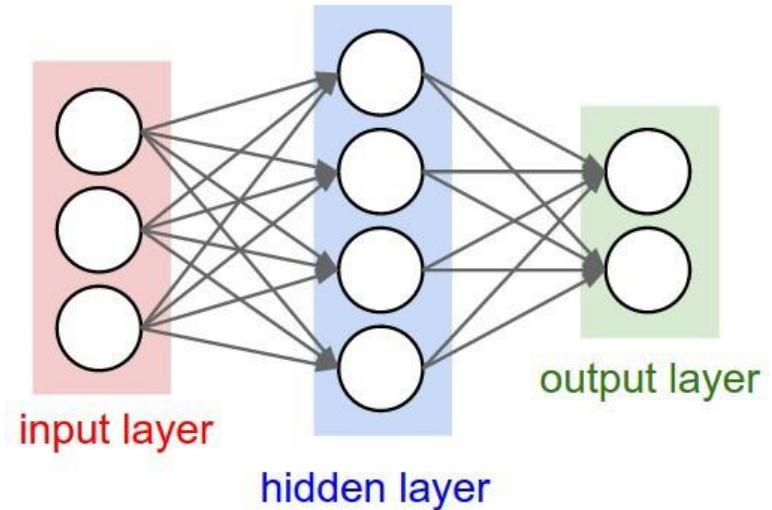
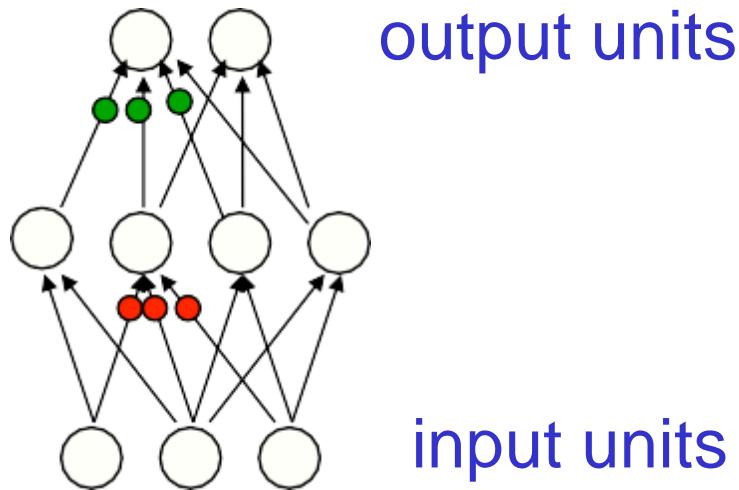
Sigmoid:  $\sigma(z) = \frac{1}{1+\exp(-z)}$

Tanh:  $\tanh(z) = \frac{\exp(z)-\exp(-z)}{\exp(z)+\exp(-z)}$

ReLU (Rectified Linear Unit):  $\text{ReLU}(z) = \max(0, z)$



Network with one layer of four hidden units:



**Figure :** Two different visualizations of a 2-layer neural network. In this example: 3 input units, 4 hidden units and 2 output units

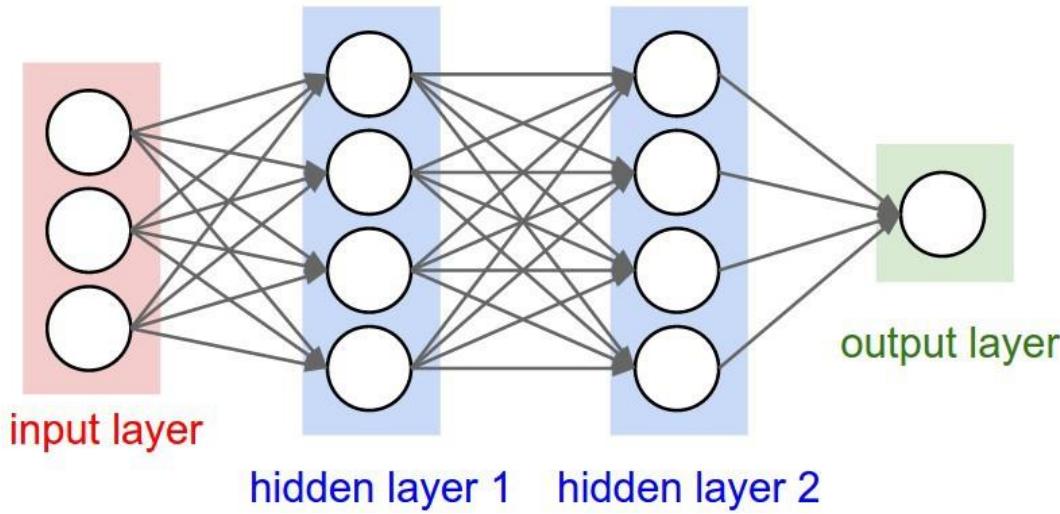
Naming conventions; a 2-layer neural network:

One layer of hidden units

One output layer

(we do not count the inputs as a layer)

Going deeper: a 3-layer neural network with two layers of hidden units



**Figure :** A 3-layer neural net with 3 input units, 4 hidden units in the first and second hidden layer and 1 output unit

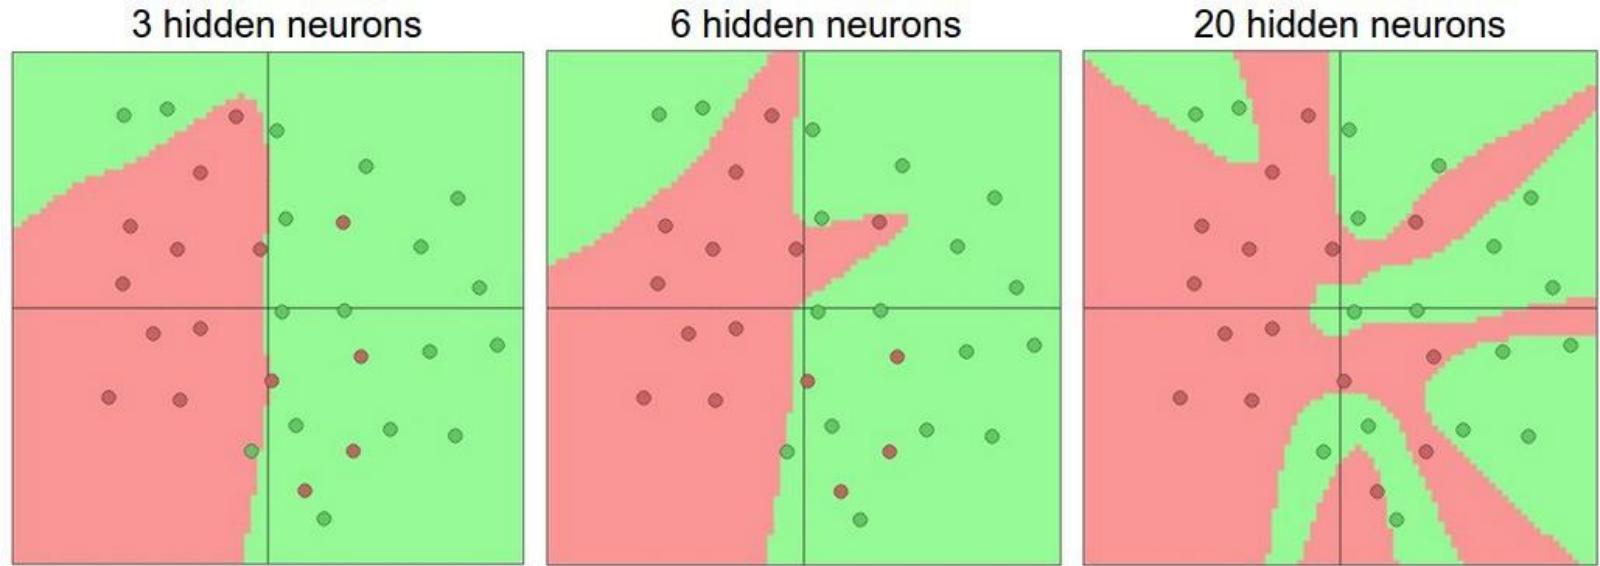
Naming conventions; a N-layer neural network:

$N - 1$  layers of hidden units

One output layer

Neural network with at **least one hidden layer** is a universal approximator (can represent any function).

Proof in: Approximation by Superpositions of Sigmoidal Function, Cybenko, [paper](#)

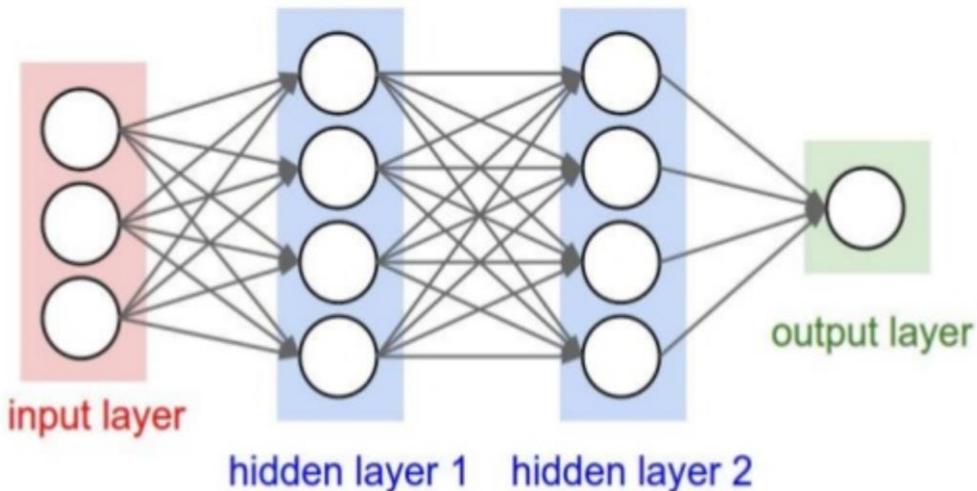


The capacity of the network increases with more hidden units and more hidden layers

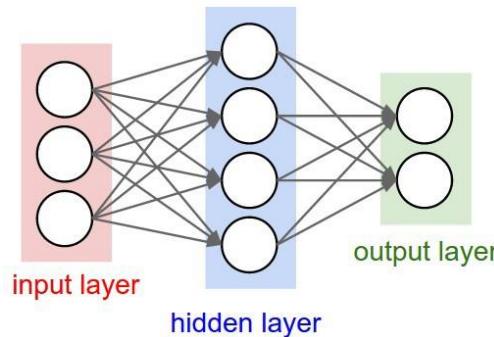
Why go deeper? Read e.g.,: Do Deep Nets Really Need to be Deep? Jimmy Ba, Rich Caruana, Paper: [paper](#)]

# Forward Pass (Forward Propagation)

- The input neurons first receive the data features of the object. After processing the data, they send their output to the first hidden layer.
- The hidden layer processes this output and sends the results to the next hidden layer.
- This continues until the data reaches the final output layer, where the output value determines the object's classification.
- This entire process is known as **Forward Propagation**, or **Forward prop**.



# Forward Pass (Forward Propagation)



Output of the network can be written as:

$$h_j(\mathbf{x}) = f(v_{j0} + \sum_{i=1}^D x_i y_i)$$

$$o_k(\mathbf{x}) = g(w_{k0} + \sum_{j=1}^J h_j(\mathbf{x}) w_k)$$

( $j$  indexing hidden units,  $k$  indexing the output units,  $D$  number of inputs)

Activation functions  $f, g$ : sigmoid/logistic, tanh, or rectified linear (ReLU)

$$\sigma(z) = \frac{1}{1 + \exp(-z)}, \quad \tanh(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)}, \quad \text{ReLU}(z) = \max(0, z)$$

# Backward Pass (Back-propagation)

- **Back-propagation**: an efficient method for computing gradients needed to perform gradient-based optimization of the weights in a multi-layer network

## Training neural nets:

Loop until convergence:

- ▶ for each example  $n$ 
  1. Given input  $\mathbf{x}^{(n)}$ , propagate activity forward ( $\mathbf{x}^{(n)} \rightarrow \mathbf{h}^{(n)} \rightarrow \mathbf{o}^{(n)}$ ) (**forward pass**)
  2. Propagate gradients backward (**backward pass**)
  3. Update each weight (via gradient descent)

# What are we propagating Backward?

- Find weights:

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_{n=1}^N \text{loss}(\mathbf{o}^{(n)}, \mathbf{t}^{(n)})$$

where  $\mathbf{o} = f(\mathbf{x}; \mathbf{w})$  is the output of a neural network

- Define a loss function, eg:

- ▶ Squared loss:  $\frac{1}{2} \sum_k (o_k^{(n)} - t_k^{(n)})^2$
- ▶ Cross-entropy loss:  $-\sum_k t_k^{(n)} \log o_k^{(n)}$

- Gradient descent:

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \frac{\partial E}{\partial \mathbf{w}^t}$$

where  $\eta$  is the learning rate (and  $E$  is error/loss)

# Some Types of Deep Neural Networks

- **Feedforward Networks (FFN):** Best for when we have an idea of the output
- **Recurrent Neural Network (RNN):** Best for time series Data
- **Convolutional Neural Network (CNN):** Best of Image, text, voice recognition.

# Deep Neural Networks Weaknesses

- Deep Learning **requires a large dataset**, hence long training period.
- In term of cost, Machine Learning methods like SVMs and other tree ensembles are very easily deployed even by relative machine learning novices and can usually get you reasonably good results.
- Deep learning methods **tend to learn everything**. It's better to encode prior knowledge about structure of images (or audio or text).
- The learned features are often **difficult to understand**. Many vision features are also not really human-understandable (e.g, concatenations/combinations of different features).
- Requires **a good understanding of how to model multiple modalities with traditional tools**.

# Useful Tools

- TensorFlow
- PyTorch
- Apache MXNet
- Caffe

# Q & A