COMP9321 Data Services Engineering

Term3, 2019

Week 10: Final Wrap-up

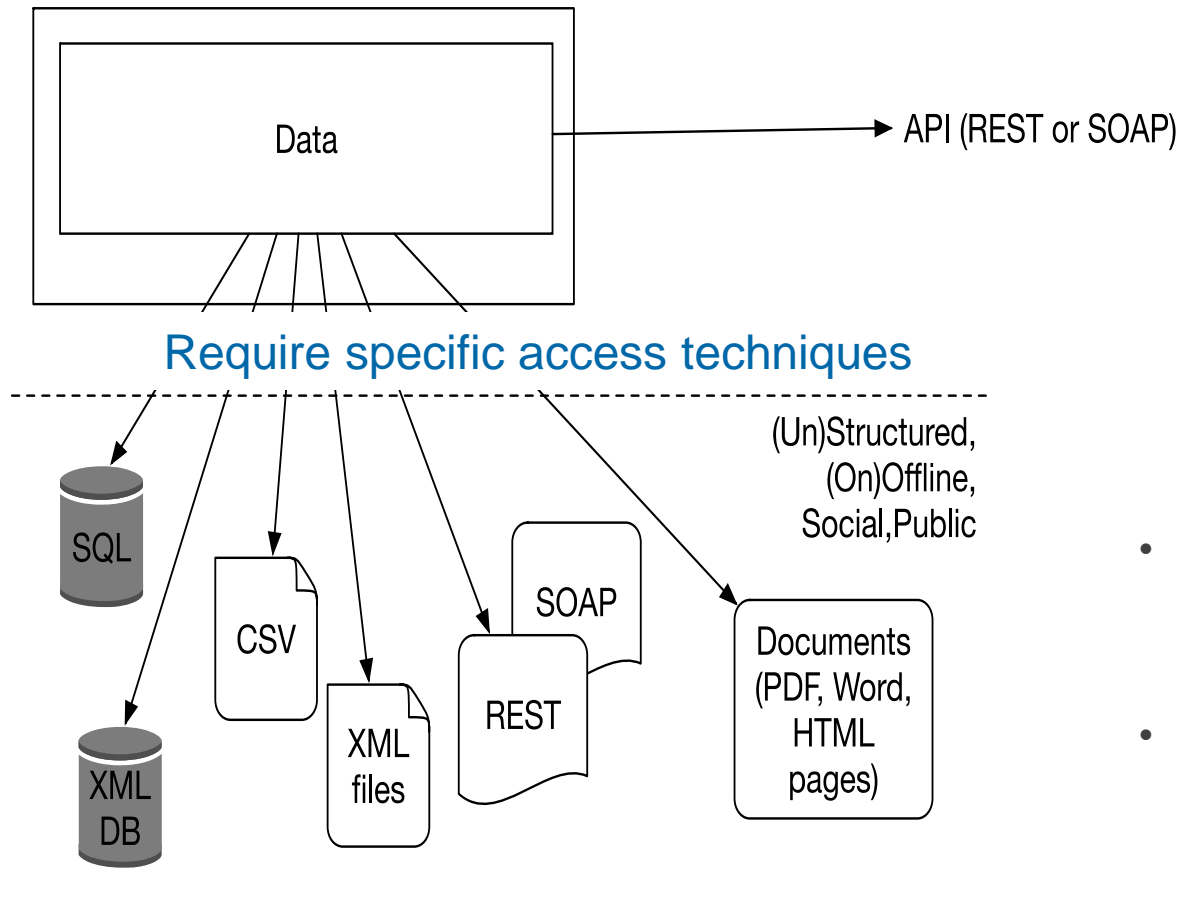**Mortada Al-Banna, CSE UNSW Australia**

Course Aims

This course aims to introduce the student to core concepts and practical skills for engineering the data in service-oriented data-driven applications. Specifically, the course aims to answer these questions:

- *How to access and ingest data from various external sources?*

- *How to process and store the data?*

- *How to curate (e.g. Extract, Transform, Correct, Aggregate, and Merge/Split) and publish the data?*

- *How to apply available analytics to the data?*

- *How to visualize the data to communicate effectively*

Data Services – what is it about?

Two sides of a coin:

Data → API (REST or SOAP)

Require specific access techniques

SQL

CSV

XML files

SOAP

REST

XML DB

(Un)Structured, (On)Offline, Social,Public

Documents (PDF, Word, HTML pages)

- Data integration/aggregation from multiple sources (data prep)

- Data publication for consumer access (API)

UNSW SYDNEY

Understanding the Data (ask the right Questions)

- What is this dataset?

- What should I expect within this dataset?

- Basic concepts (e.g., domain knowledge)

- What are the questions that I need to answer?

- Does the dataset have some sort of a schema? (utilize domain knowledge)

Obtaining Data

Useful data can be found in many places

• on the Web, possibly via an API

• in documents in a file system

• in spreadsheets

• in videos….etc. etc. etc.

and in a variety of formats

• Unformatted text (in files)

• PDF documents (in files)

• HTML documents (web pages)

• XML documents (via web APIs)

• JSON data (often via web APIs)

• CSV data files (spreadsheets)

Relational Model vs. "NoSQL" Models

Relational Model (more or so synonymous with SQL)

- The best known, probably the most successful data model which has proven itself in many aspects to be the data model of choice in many applications

- Data is organised into relations (table) where each relation holds an unordered collection of tuples (rows)
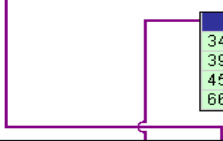
Based on solid theory and well engineered implementation -> many competing models have been proposed, but never managed to take over SQL

Built for business data processing

- Typical business transactions (airline reservations, stock keeping, etc.)
- Batch processing (invoicing, payroll, reporting, etc.)

Turned out it was still generically applicable to many modern Web applications too

Hypothetical Relational Database Model

PubID	Publisher	PubAddress
03-4472822	Random House	123 4th Stree, New York
04-7733903	Wiley and Sons	45 Lincoln Blvd, Chicago
03-4859223	O'Reilly Press	77 Boston Ave, Cambridge
03-3920886	City Lights Books	99 Market, San Francisco

AuthorID	AuthorName	AuthorBDay
345-28-2938	Haile Selassie	14-Aug-92
392-48-9965	Joe Blow	14-Mar-15
454-22-4012	Sally Hemmings	12-Sep-70
663-59-1254	Hannah Arendt	12-Mar-06

| ISBN	AuthorID	PubID	Date	Title
1-34532-482-1	345-28-2938	03-4472822	1990	Cold Fusion for Dummies
1-38482-995-1	392-48-9965	04-7733903	1985	Macrame and Straw Tying
2-35921-499-4	454-22-4012	03-4859223	1852	Fluid Dynamics of Aquaducts
1-38278-293-4	663-59-1254	03-3920886	1967	Beads, Baskets & Revolution

UNSW SYDNEY

Relational Model vs. "NoSQL" Models

The rise of NoSQL … (since 2010 or so)

• Refers to a host of technologies that implement distributed, "non-relational" databases

Why NoSQL?

• A need for greater scalability – very large datasets or very high 'write' throughput

• A need for more expressive and dynamic data model

• Usually do not require a fixed table schema nor do they use the concept of joins

• All NoSQL offerings relax one or more of the ACID properties

Data Cleansing

- Dealing with Missing Data

- Removing Unnecessary Data (rows, or columns)

- Normalizing/Formatting data

Manipulating the data

- Merging Data

- Applying a function to data

- Pivot tables

- Change the index of a dataframe

- Groupby

Data Visualisation

Referring to any visual representation of data that is:

- algorithmically drawn (may have custom touches but is largely rendered with the help of computerized methods);

- easy to regenerate with different data (the same form may be repurposed to represent different datasets with similar dimensions or characteristics);

- often aesthetically simple (data is not decorated); and

- relatively data-rich (large volumes of data are welcome and viable, in contrast to infographics).

UNSW
SYDNEY

What You Need to Cater for Data Visualisation
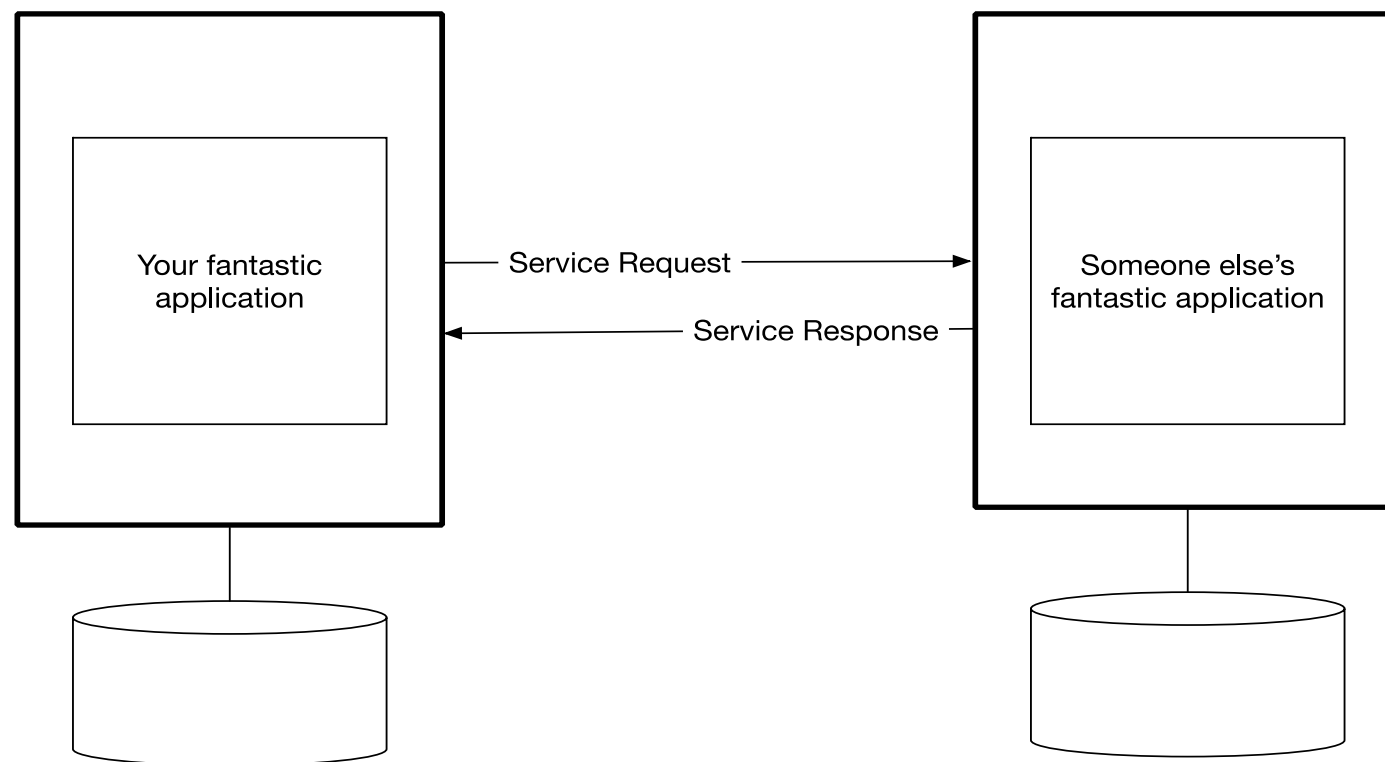
- Accuracy is important, having a clear story to tell is important

- You need to be ready to do some basic data prep and pre analysis before visualisation

- Knowing the right paradigm (form) to use for the story

- Aware of your own limitation as 'non-expert' (visualisation is not easy)

Actually, a lot of experts recommend "sketching the idea out" with pen and paper.

API – Application Programming Interface

Your fantastic application → Service Request → Someone else's fantastic application

Someone else's fantastic application → Service Response → Your fantastic application

- The interface is not meant for human interactions – there is another program on the other side → implication of this: you must have a clear contract (e.g., IOU Alice Bob 100)

- These days companies use APIs internally (private APIs) as well as exposing them externally (public APIs)

What is and Why a RESTful Service

REST is an architectural style of building networked systems - a set of architectural constraints in a protocol built in that style.

The protocol in REST is HTTP (the core technology that drives the Web)

Popular form of API ... It is popularised as a guide to build modern distributed applications on the Web – let's work with the components that the Web itself is built in.

REST itself is not an official standard specification or even a recommendation. It is just a "design guideline" for building a system (or a service in our context) on the Web

Architectural Constraints of REST

1. Client-Server

2. Uniform Interface

3. Statelessness

4. Caching

5. Layered System

6. Code on demand (optional)

Designing RESTful APIs

A well-designed API should make it easy for the clients to understand your service without having to "study" the API documents in-depth.

self-describing, self-documenting as much as possible

the clients are developers like yourself, so probably they would like to have an API that is easy to pick up and go

The RESTful service principles actually give us a straightforward guideline for designing the Web API ...

"Clean, Clear, Consistent" are the key

REST API Security

- REST API security is Important, Why?

- It matter enough that OWASP included many instances in their web security Top ten related to APIs and they have the REST Security cheat sheet.

- REST relies on the elements of the Web for security too (Check OWASP top 10)

- HTTPS (SSL) "Strong" server authentication, confidentiality and integrity protection The only feasible way to secure against man-in-the-middle attacks

- Any security sensitive information in REST API should use SSL

- Other things to remember (Input Validation, Methods restriction, logging)

https://www.owasp.org/images/7/72/OWASP_Top_10-2017_%28en%29.pdf.pdf
https://www.owasp.org/index.php/REST_Security_Cheat_Sheet

UNSW
SYDNEY

REST APIs and Security (Cont'd)

API developers at least must deal with authentication and authorisation: Authentication (401 Unauthorized) vs. Authorisation (403 Forbidden):

Common API authentication options:

- HTTP Basic (and Digest) Authentication: IETF RFC 2617
- Token-based Authentication
- API Key [+ Signature]
- OAuth (Open Authorisation) Protocol - strictly uses HTTP protocol elements only

https://www.owasp.org/index.php/Transport_Layer_Protection_Cheat_Sheet
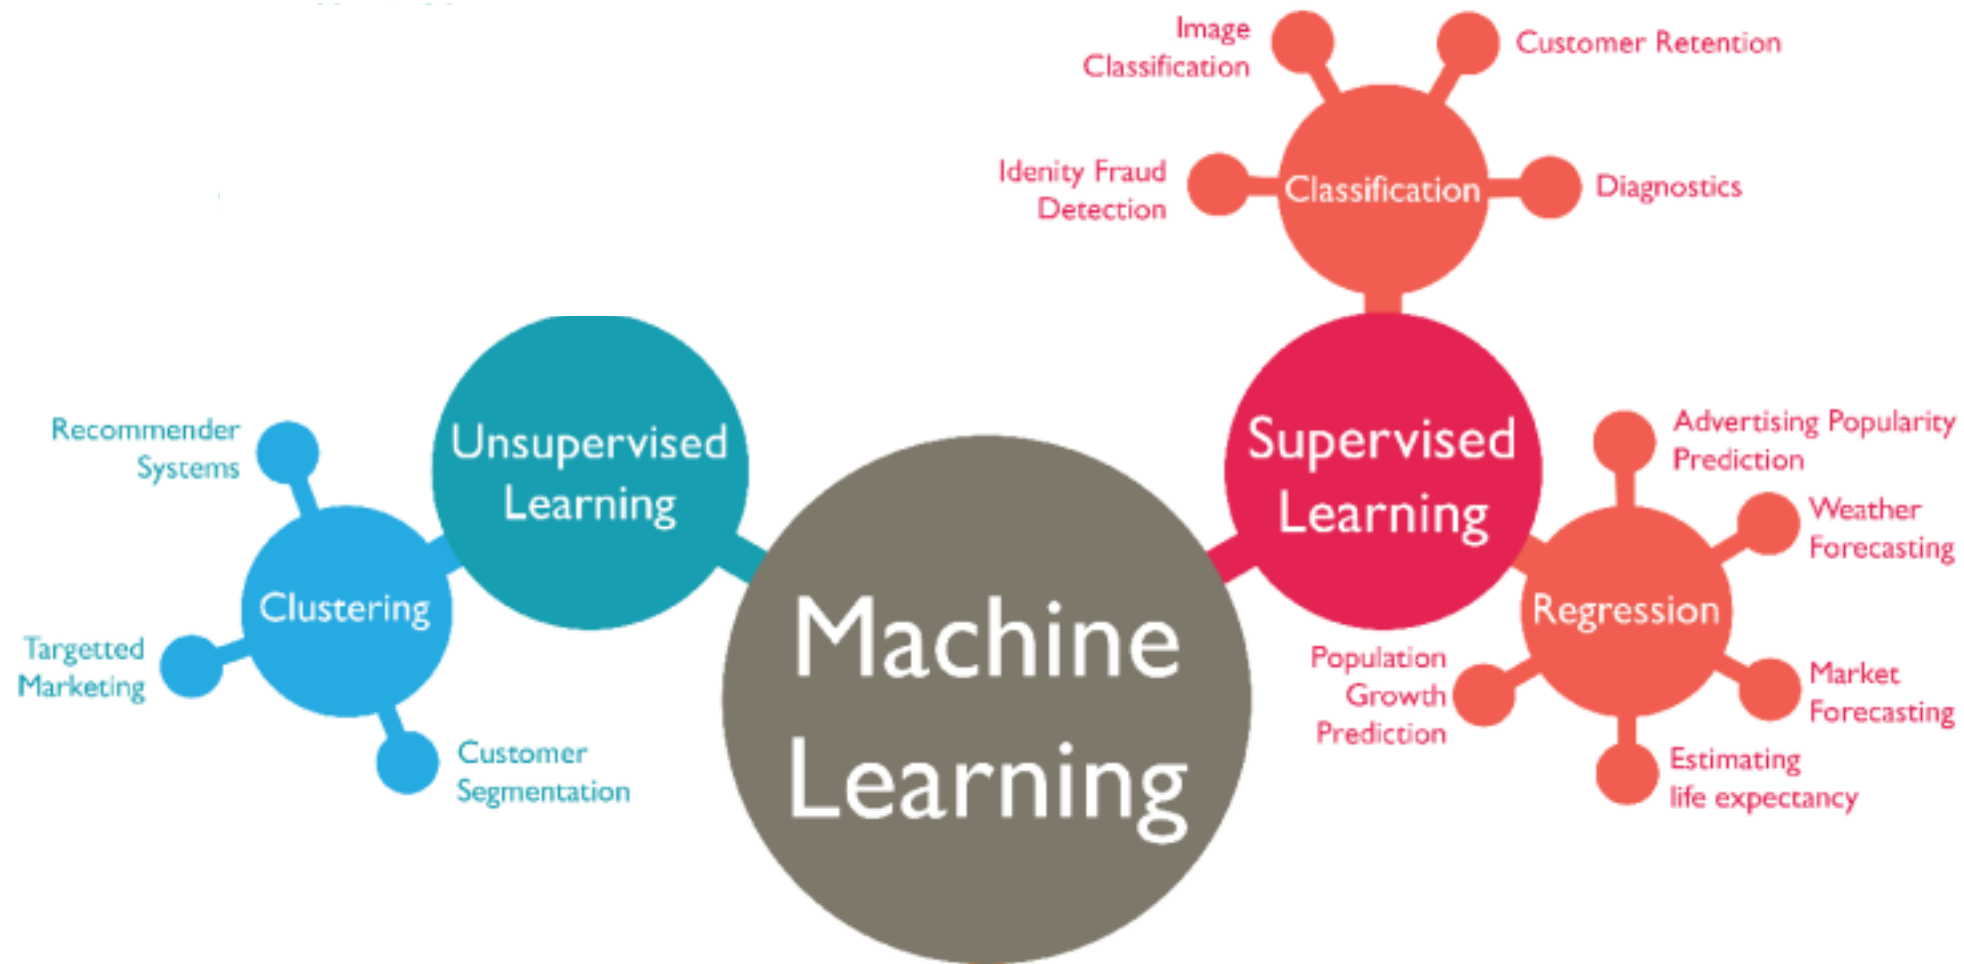
RESTful Service Client
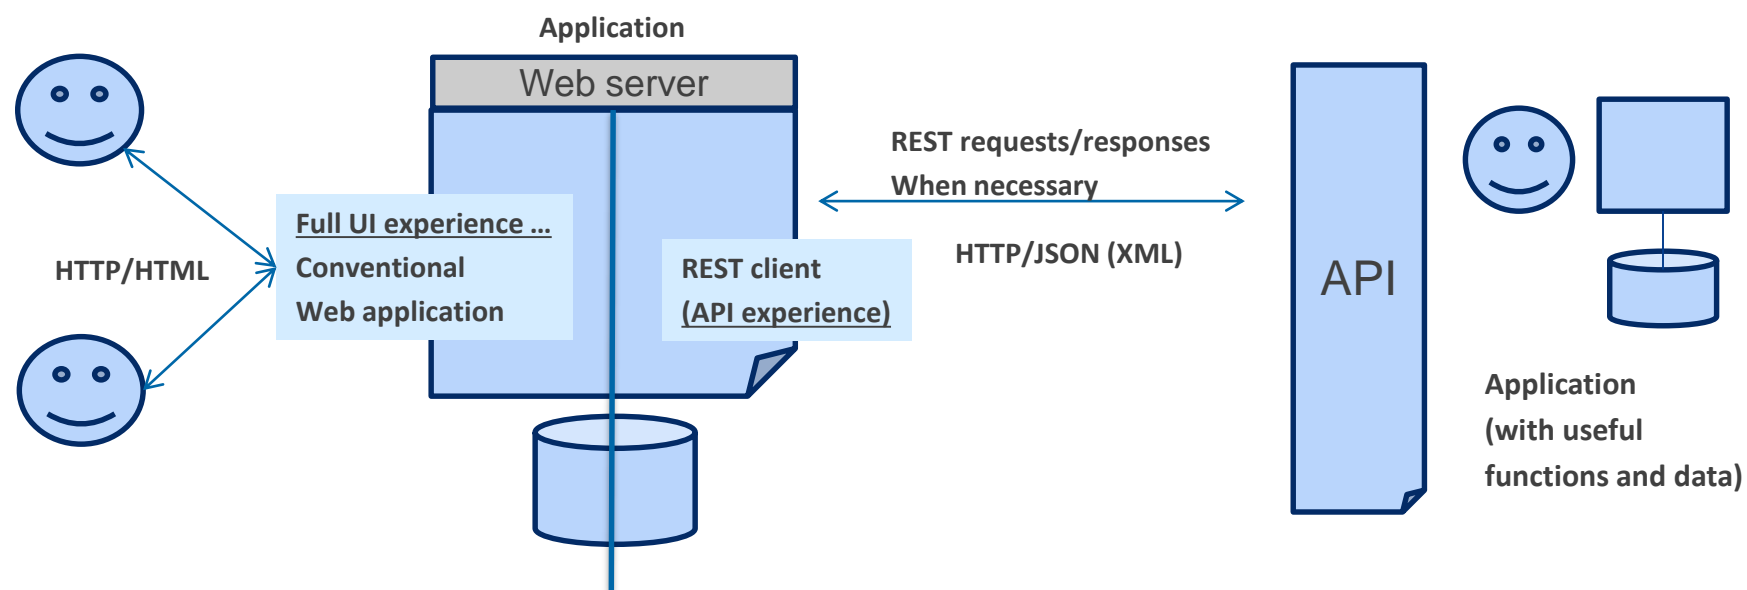
Machine Learning for Data Analytics

Machine Learning for Data Analytics

1. **Define** and **Initialize** a Model

2. **Train** your Model (using your training dataset)

3. **Validate** the Model (by prediction using your test dataset)

4. Use it: **Explore** or **Deploy** as a web service

5. **Update** and **Revalidate**

UNSW
SYDNEY

https://www.datasciencecentral.com/profiles/blogs/decoding-machine-learning-methods

Recommender Systems

Given:

- User model (e.g. ratings, preferences, demographics, situational context)

- Items (with or without description of item characteristics)

Find:

- Relevance score. Used for ranking.

Deep Learning

Neural networks define functions of the inputs (hidden features), computed by neurons

Artificial neurons are called units

Figure : A mathematical model of the neuron in a neural network

Wrap-up Example

- **Task:** I want to find candidate security researchers to discover vulnerabilities in my systems

- **Possible Data sets**:
  - Hackerone full disclosure
  - OpenBugBounty
  - ExploitDB

Example Cont'd

- Working with the Data
  - Understand the Data: visit the platforms offering the data and start browsing or inspecting the few records. Understand what is the type of the data and the schema and the meta-data.
  - Acquire the Data: Download the data file, query the data source, or dump the database if possible. Store the data in your selected mode of storage.
  - Clean the Data: drop the useless fields or the ones with missing or corrupted data.
  - Augment the Data: it might be useful to combine two data sources to acquire better insight.

Example Cont'd

- Design the Solution
  - What is the problem that you need to solve. What are the features you need to provide.
  - What kind of Data analytics you need to perform
    - Visualize your Data
    - What Machine learning model
    - What are your constrains
  - Design the API endpoints to allow the consumption of your service
  - Document everything

Assessment

Assessment:

- 40% formal written exam: individual assessment.
- 50% on assignment work
  - Assignment1 on Working with Data(individual) 20%
  - Assignment2 on building a data analytics service (group) 30%
- 10% on 5 online quizzes (WebCMS-based quiz system, 'open' test)

Final Mark = quizzes + assignments + exam

The Final Exam

- Duration: Two hours

- Lab Based Exam

- Types of questions:
    - Multiple Choice
    - Short Answers
    - Written Answers

The Final Exam (examples of Questions)

- Multiple Choice Question:

In REST Services, Stateless means that every HTTP request

○	have a fixed state
○	happens in a complete isolation.
○	connected to other requests
○	can not have a certain state

The Final Exam (examples of Questions)

- Short Answer:

Many kinds of data can be modelled as a graph. For example, in a **(blank A)** graph, the vertices represent people and edges indicate which people know each other. In a **(blank B)** graph, vertices are web pages and edges indicate HTML links to other pages.

Write down the word(s) that could fill in **blank A** and **blank B**.

The Final Exam (examples of Questions)

- Written Answer:

Question: Explain the safeness and idempotence properties of REST principles.

The Final Exam (examples of Questions)

- Consider the following HTTP request invoking a POST method of a RESTful API:

POST /orders HTTP/1.1

Host: api.coffeehouse.com

Content-Type: application/xml

<order>

<drink>latte</drink>

</order>

Write down the content of the HTTP response that you would return as the result

Supplementary Exam Policy

Supp Exam is only available to students who:

• DID NOT attend the final exam

• Have a good excuse for not attending

• Have documentation for the excuse

Submit special consideration within 72 hours (via myUNSW with supporting docs)

Everybody gets exactly one chance to pass the final exam. For CSE supplementary assessment policy, follow the link in the course outline.

Warning…
Warning

Huge Thanks Coming Your Way…

UNSW
SYDNEY

Thanks for the Teaching Team

- Course Administrator
  - Mohammad Ali Yaghub Zade Fard

- Tutors
  - Mohammad Ali Yaghub Zade Fard
  - May Altulyan
  - Rajitha Yasaweerasinghelage
  - Madhushi Bandara

That's all Folks!